




Aliasing: Risks, Opportunities and Techniques

Roi Barkan

[Slides](https://www.youtube.com/watch?v=1eAERikCzVg)

<https://www.youtube.com/watch?v=1eAERikCzVg>



*“All problems in computer science can be solved by another level of indirection”
“... except for the problem of too many levels of indirection”*

- David Wheeler

levelofindirection.com

Hi, I'm Roi

- Roi Barkan (he/him) - רועי ברקן
- I live in Tel Aviv, Israel
- C++ developer since 2000
- VP Technologies @ Istra Research
 - Finance, Low Latency, in Israel
 - careers@istraresearch.com
- My first on site international C++ conference
 - Please - ask questions, make comments

What is Aliasing?

- Definition: two (or more) variables which refer to the same memory location being used together.
- Example:

```
std::string s{"hello, "};  
s += s;
```
- Causes dependencies to exist where the code seems independant
 - Aliasing is NOT about threads and volatile data
 - Reasoning about aliasing can be similar to reasoning about race conditions
- Aliasing considerations impact code correctness and efficiency/speed

Talk Outline

- Examples
 - Correctness, Performance
- Aliasing and the C++ Standard
- Dealing with aliasing pitfalls
 - APIs and implementations
 - Standard vs. compiler specific
- Future of aliasing
- Aliasing based design

Example: Aliased Function Arguments

- Pointers:

```
auto minmax = [](const string& i, const string& j,  
                 string* out_min, string* out_max) {  
    *out_min = min(i, j); *out_max = max(i, j);  
};  
array<string, 2> arr{"22222", "11111"};  
minmax(arr[0], arr[1], &arr[0], &arr[1]); // try to sort
```

- References:

```
auto concat = [](string& result, const auto&... args) {  
    ((result += args), ...);  
};  
string x{"hello "}, y{"world "};  
concat(x, y, x);
```

```

10 int main() {
11     auto minmax = [](const string& i, const string& j, string* out_min,
12                     string* out_max) {
13         *out_min = min(i, j);
14         *out_max = max(i, j);
15     };
16     array<string, 2> arr{"22222", "11111"};
17     // try to sort
18     minmax(arr[0], arr[1], &arr[0], &arr[1]);
19     cout << "expect 22222 and get " << arr[1] << "\n";
20     auto concat = [](string& result, const auto&... args) {
21         ((result += args), ...);
22     };
23     string x{"hello "}, y{"world "};
24     concat(x, y, x);
25     cout << "expect [hello world hello ] and get [" << x << "]\n";
26     return 0;


```


Executor x86-64 clang 14.0.0 (C++, Editor #1) ✎ ✕

A ▾

☐ Wrap lines

 Libraries (1)

 Compilation

 Arguments

 Stdin

 Compiler output

x86-64 clang 14.0.0 ▾



-std=c++20 -O3

Program returned: 0

Program stdout

expect 22222 and get 11111

expect [hello world hello] and get [hello world hello world]

Example: Not Only Arguments

- Member variables:

```
complex<int> x{2, 2};  
x *= reinterpret_cast<int*>(&x)[0]; // multiply by real part
```

- Lambda closures:

```
auto add_to_all = [](auto& v, const auto& val) {  
    for_each(begin(v), end(v), [&](auto& x) { x += val; });  
};  
vector<int> v{1, 2, 3};  
add_to_all(v, v[0]);
```




```

12 //    members();
13     complex<int> x{2, 2};
14     x *= reinterpret_cast<int*>(&x)[0]; // multiply by real part
15     cout << "expect (4,4) and get " << x << "\n";
16 //    lambdas();
17     auto add_to_all = [](auto& v, const auto& suffix) {
18         for_each(begin(v), end(v), [&](auto& x) { x += suffix; });
19     };
20     vector<int> v{1, 2, 3};
21     add_to_all(v, v[0]);
22     cout << "expected [2,3,4] and got [" << v[0] << "," << v[1] << "," << v[2]
23         << "]\n";

```

Executor x86-64 clang 14.0.0 (C++, Editor #1)  

A ☐ Wrap lines **Libraries (1)**  Compilation **>** Arguments **↔** Stdin **↔** Compiler output

x86-64 clang 14.0.0 ▼



-std=c++20 -O3

Program returned: 0

Program stdout

expect (4,4) and get (4,8)

expected [2,3,4] and got [2,4,5]

Example: Aliased Buffers

```
void loopcpy(char* dst, const char* src, int size) {  
    while (size-->0) *dst++ = *src++;  
}
```

```
test("loopcpy", loopcpy);  
test("strcpy ", [](auto dst, auto src, auto size) {  
    strcpy(dst, src); });  
test("strncpy", strncpy);  
test("memcpy ", memcpy);  
test("memmove", memmove);  
test("copy_n ",  
    [](auto dst, auto src, auto size) {  
        copy_n(src, size, dst); });
```

Clang14

loopcpy [hhhhhh] Bad

strcpy [helll] Bad

strncpy [hello] Good

memcpy [hello] Good

memmove [hello] Good

copy_n [hello] Good

ICC 2021.5.0

[hhhhhh] Bad

[helll] Bad

[hello] Good

[helll] Bad

[hello] Good

[hello] Good

Standard

Bad

UB

UB

UB

Good

ID

Example: STL Algorithms

- Erase (or Erase-Remove) max element with duplicates

```
erase(v, *max_element(begin(v), end(v)));
```

or (C++20 ranges)

```
erase(v, *ranges::max_element(v));
```

- (**remove** has [documentation](#) about this, **erase** doesn't)

- Copy/Move overlapping regions

```
copy(begin(v), end(v)-1, begin(v)+1);
```

- ([Documented](#) as faulty, **copy_backward** recommended instead)

- Iterators can cause aliasing

```
auto max = ranges::max_element(a);
```

```
stable_partition(begin(a), end(a), [=] (const auto&x) {return x != *max;});
```

- (Predicates which modify their argument or the sequence are UB, this case isn't)

```

7 void erase() {
8     vector<int> v{1, 4, 2, 1, 4, 3, 4};
9     erase(v, *max_element(begin(v), end(v)));
10    cout << "erase_max          expected [1,2,1,3]      and got [";
11    copy(begin(v), end(v) - 1, ostream_iterator<int>(cout, ","));
12    cout << v.back() << "]\n";
13    erase(v, *ranges::max_element(v));
14    cout << "erase_ranges::max expected [1,2,1,3]      and got [";
15    copy(begin(v), end(v) - 1, ostream_iterator<int>(cout, ","));
16    cout << v.back() << "]\n";
17 }

```

Executor x86-64 clang 14.0.0 (C++, Editor #1) ✎ ✕

A ▾
☐ Wrap lines
 Libraries (1)
⚙️ Compilation
>_ Arguments
↔ Stdin
↔ Compiler output

x86-64 clang 14.0.0 ▾



-std=c++20 -O3

Program returned: 0

Program stdout

```

erase_max          expected [1,2,1,3]      and got [1,2,1,4,3,4]
erase_ranges::max  expected [1,2,1,3]      and got [1,2,1,3,4]
copy               expected [a,b,c,d]      and got [a,b,b,b]
stable_partition   expected [1,2,3,4,4,4]  and got [1,2,4,3,4,4]

```

```

19 void copy() {
20     vector<string> v{"b", "c", "d", "e"};
21     copy(begin(v), end(v) - 1, begin(v) + 1);
22     v[0] = "a";
23     cout << "copy          expected [a,b,c,d]    and got [";
24     copy(begin(v), end(v) - 1, ostream_iterator<string>(cout, ","));
25     cout << v.back() << "]\n";
26 }
27
28 void partition() {
29     array a = {1, 4, 2, 4, 3, 4};
30     auto max = ranges::max_element(a);
31     stable_partition(begin(a), end(a),
32         [=](const auto& x) { return x != *max; });
33     cout << "stable_partition expected [1,2,3,4,4,4] and got [";
34     copy(begin(a), end(a) - 1, ostream_iterator<int>(cout, ","));
35     cout << a.back() << "]\n";

```

Executor x86-64 clang 14.0.0 (C++, Editor #1) ✎ ✕

A ▾

☐ Wrap lines

 Libraries (1)

 Compilation

▸ Arguments

↔ Stdin

 Compiler output

x86-64 clang 14.0.0 ▾



-std=c++20 -O3

Program returned: 0

Program stdout

copy expected [a,b,c,d] and got [a,b,b,b]

stable_partition expected [1,2,3,4,4,4] and got [1,2,4,3,4,4]

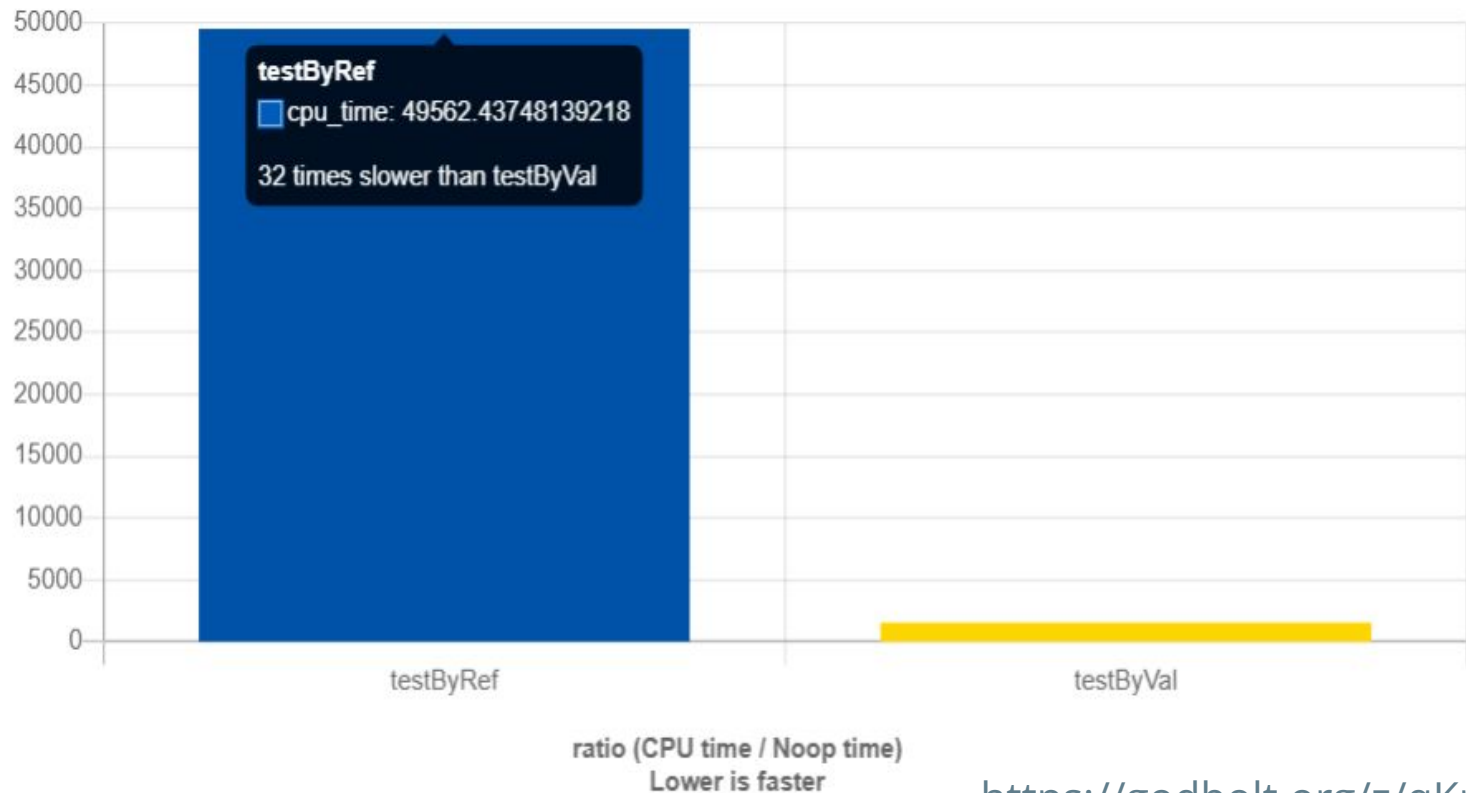
Performance Effect of Aliasing

- Extreme example

```
void foo(std::vector<double>& v, const double& coeff) {  
    for (auto& item : v) item *= std::sin(coeff);  
}
```

- Compiler's missed opportunities:
 - Register <-> memory
 - Vectorization
 - Expression hoisting
- How important can it be...

Performance Benchmark Results



Lesson Learned - Aliasing is Tricky

- Humans rarely consider it → Strange unexpected bugs
 - We expect independence of different variables
- Compilers can't ignore it → Unexpected performance loss
 - Learn more in Ofek Shilon's talk about optview2 on Wednesday
- Library writers should document it → users should read documentation
 - Misuse often leads to 'happens to work' code
- *“All problems in computer science can be solved by another level of indirection”*
“... except for the problem of too many levels of indirection”

Aliasing in Other Languages

- The C language had a surge of (non-assembly) aliasing issues
 - Pointers were used as primitive substitutes to arrays, matrices, strings
 - C99 introduced the `restrict` keyword
 - A code block with a `restrict` pointer/array can only change the pointed data through that pointer/array. Otherwise: undefined behavior
 - Most C++ compilers have some non-standard support for `restrict`
- Fortran typically treats aliases as undefined behavior
 - with compiler switches to assume aliasing
- Swift and Rust track reference creation aiming to prohibit the risk of aliasing

Aliasing in the C++ Standard

- C++ hasn't adopted the **restrict** keyword (yet?)
 - Seems more tricky: function-signature qualifiers, templates, functors/lambda's
- Aliasing *should* be type-based - known as "strict aliasing"
 - Only similar types are technically allowed to alias each other (and **char**, **std::byte**)
 - Similar types - changes to const/volatile/signed, or base-derived relationship
 - Otherwise - undefined behavior
 - Strong-typedefs can reduce risk and improve performance !
 - Most compiler-optimizers relax the rules - favoring predictability over performance
 - Still - compilers try to prove whether aliasing is impossible
- The STL tries to document the effect of aliasing and sometimes mitigates
 - `vec.push_back(v.front())` ; always works (with a performance cost)
 - `std::bind()` holds its 'closure' by-value and avoids aliasing

Strong Typedefs

- Types that encapsulate and behave like other types, but are different and don't automatically convert to/from them
 - No standard implementation, but a few libraries mimic the behavior

- Motivating example:

```
struct A { int i; };  
struct B { int i; };
```

```
int mayAlias(auto& a, const auto& b) {  
    a.i += b.i;  
    if (b.i == 2) return 0;  
    return 1;  
}
```

```
template int mayAlias(A&, const A&);  
template int mayAlias(A&, const B&);
```

```
int mayAlias<A, A>(A&, A const&):  
    mov     eax, DWORD PTR [rsi]  
    add     DWORD PTR [rdi], eax  
  
    xor     eax, eax  
    cmp     DWORD PTR [rsi], 2  
    setne   al  
    ret  
  
int mayAlias<A, B>(A&, B const&):  
    mov     eax, DWORD PTR [rsi]  
    add     DWORD PTR [rdi], eax  
  
    cmp     eax, 2  
    setne   al  
    movzx   eax, al  
    ret
```

How to Avoid Aliasing Pitfalls

- Pass arguments by value
 - Value semantics are all the rage
 - Move semantics and copy-elision can make this relatively cheap
 - Consider supporting `std::reference_wrapper` (i.e. `std::ref()`)
- Use strong typedefs and unit libraries
 - clearer code for humans, compilers might optimize it as well
- Document your code's aliasing assumptions (contract)
 - Read other people's documentation
- For a large user base - write defensive code
 - Verify your contract - assert/throw/etc.
 - Widen your contract (e.g. `vec.push_back(v.front())`)
 - Let users control the contract

Defensive Code

- Basic function

```
template <typename Value, typename BinOp>
void unsafe_apply(std::span<Value> s, const Value& v, BinOp op) {
    for (auto& item : s) item = op(item, v);
}
```

- User controlled version

```
template <typename T> struct ByRef { using type = const T&; };
template <typename T> struct ByVal { using type = T; };

template <typename Value, typename BinOp, typename PassBy = ByRef<Value>>
void user_apply(std::span<Value> s, const Value& v_ref, BinOp op, PassBy = {}) {
    typename PassBy::type v{v_ref};
    for (auto& item : s) item = op(item, v);
}
```

Defensive Code

- Safe version

```
template <typename Value, typename BinOp>
void safe_apply(std::span<Value> s, const Value& v, BinOp op) {
    if (!s.empty() && std::less_equal{}(&s.front(), &v) &&
        std::less_equal{}(&v, &s.back()))
    {
        user_apply(s, v, op, ByVal<Value>{});
        return;
    }
    user_apply(s, v, op, ByRef<Value>{});
}
```

- Sometimes bounds/alias checking isn't as easy

Proposals on Aliasing in C++

- The **restrict** keyword signal to users and compiler that aliasing is UB
 - Many compilers have some support for it, but standardization isn't likely
- [\[\[alias_set\]\]](#) (2014) - annotate the relationship between variables
 - has some similarities with Rust lifetime annotations
- [span<T, std::restrict_access>](#) (2018) - property-based 'qualifier' for added semantics
- [std::disjoint](#) (2018) - meant for *contracts* to convey aliasing consistently
- [Lifetime safety](#) (2019) - Core guidelines and static analysis which "default to banning passing non-owning Pointers that alias".

Tricking the Compiler ?

- **union** is a mechanism for several object types to reside in the same address.
- At any time one type is *active* and accessing a different type is *typically* UB
 - **variant** is a type safe STL class that enforces correct access
- C++ does allow some accesses to non-active types - and aliasing
 - Types need to be StandardLayoutType and accessed members need to be in their common prefix. **std::is_corresponding_member** checks for this condition.
- This implies that “strict aliasing” has limits
 - I might be wrong, or this might be a bug in the standard/compilers

Strict Aliasing and **union**

x86-64 gcc 12.1



-std=c++20 -O2

A Output... Filter... Libraries + Add new...

```
1 aliasA(A&):
2     mov     eax, DWORD PTR [rdi]
3     cmp     eax, 1
4     lea     edx, [rax+rax]
5     setne   al
6     mov     DWORD PTR [rdi], edx
7     movzx   eax, al
8     ret
9 aliasU(U&):
10    mov     eax, DWORD PTR [rdi]
11    cmp     eax, 2
12    lea     edx, [rax+rax]
13    setne   al
14    mov     DWORD PTR [rdi], edx
15    movzx   eax, al
16    ret
```

- Let's add unions:

```
union U {
    A a;
    B b;
};

int aliasA(A& a) {
    return mayAlias(a, a);
};

int aliasU(U& u) {
    return mayAlias(u.a, u.b);
};
```

Different Optimizers, Different Worlds

```
x86-64 gcc 12.1  -std=c++20 -O2

9 aliasU(U&):
10     mov     eax, DWORD PTR [rdi]
11     cmp     eax, 2
12     lea     edx, [rax+rax]
13     setne   al
14     mov     DWORD PTR [rdi], edx
15     movzx   eax, al
16     ret
```

Output (0/0) x86-64 gcc 12.1 - 757ms (105298) ~700 lines filtered

x86-64 gcc 12.1 (C++, Editor #1, Compiler #1)

```
x86-64 gcc 12.1  -std=c++20 -O1

9 aliasU(U&):
10     mov     eax, DWORD PTR [rdi]
11     lea     edx, [rax+rax]
12     mov     DWORD PTR [rdi], edx
13     cmp     eax, 1
14     setne   al
15     movzx   eax, al
16     ret
```

```
x86-64 clang 14.0.0  -std=c++20 -O2

9 aliasU(U&):
10     mov     ecx, dword ptr [rdi]
11     lea     eax, [rcx + rcx]
12     mov     dword ptr [rdi], eax
13     xor     eax, eax
14     cmp     ecx, 2
15     setne   al
16     ret
```

Output (0/0) x86-64 clang 14.0.0 - cached (257708) ~480 lines

x86-64 clang 14.0.0 (C++, Editor #1, Compiler #4)

```
x86-64 clang 14.0.0  -std=c++20 -O1

9 aliasU(U&):
10     mov     ecx, dword ptr [rdi]
11     lea     eax, [rcx + rcx]
12     mov     dword ptr [rdi], eax
13     xor     eax, eax
14     cmp     ecx, 2
15     setne   al
16     ret
```

variant State Machines

- State machine is a typical case for using **variant**
 - At any point only one state is valid
- Changing the state to **T** is done via **operator=(T&&)** or **emplace<T>()**
- Different states commonly share information
 - **variant<WorkingPerson, RestingPerson>** - both states typically have a name, might inherit from **Person**.
 - Semantic strong typedefs might be identical in structure, e.g. **variant<Cat, HappyCat>**
- Sadly, state changing functions aren't allowed (UB) to read the previous state (especially relevant for **emplace<T>()**)
 - Previous state gets destructed before the new state constructor is invoked
 - STL chose performance over safety (unlike most containers)

variant State Changes

- Undefined/unexpected behavior:

```
variant<filesystem::path, string> v{"some_long_file_name"s};  
v = std::move(v); //Bad on non-variants as well  
v.emplace<filesystem::path>(std::move(get<string>(v)));  
v.emplace<filesystem::path>(get<string>(v));
```

- The proper (no copy) way is to use temporaries, and rely on move semantics:

```
v.emplace<filesystem::path>(  
    string{std::move(get<string>(v))})
```

Summary

- Aliasing is tricky - people assume independence
- Value semantics makes life simpler
- Strong typedefs can assist
- Implement and document your code with care
- Smart people in the committee are working on improvements
- Know how to communicate with others and the compiler

Thank You !!

- Happy coding !
- Questions/comments are welcome

References / Acknowledgements

- OptView2 - <https://youtu.be/nVc439dnMTk>
- [\[\[alias_set\]\]](https://wg21.link/n3988) - <https://wg21.link/n3988>
- [span<T, std::restrict_access>](https://wg21.link/p0856) - <https://wg21.link/p0856>
- [std::disjoint](https://wg21.link/p1296) - <https://wg21.link/p1296>
- [Lifetime safety](https://wg21.link/p1179) - <https://wg21.link/p1179>