

~~Everything~~ I Wish They Told Some Stuff Me About Linkers



Ofek Shilon

And
Linux/Windows
diffs
Focusing on
shared libs

And Loaders





Ofek Shilon

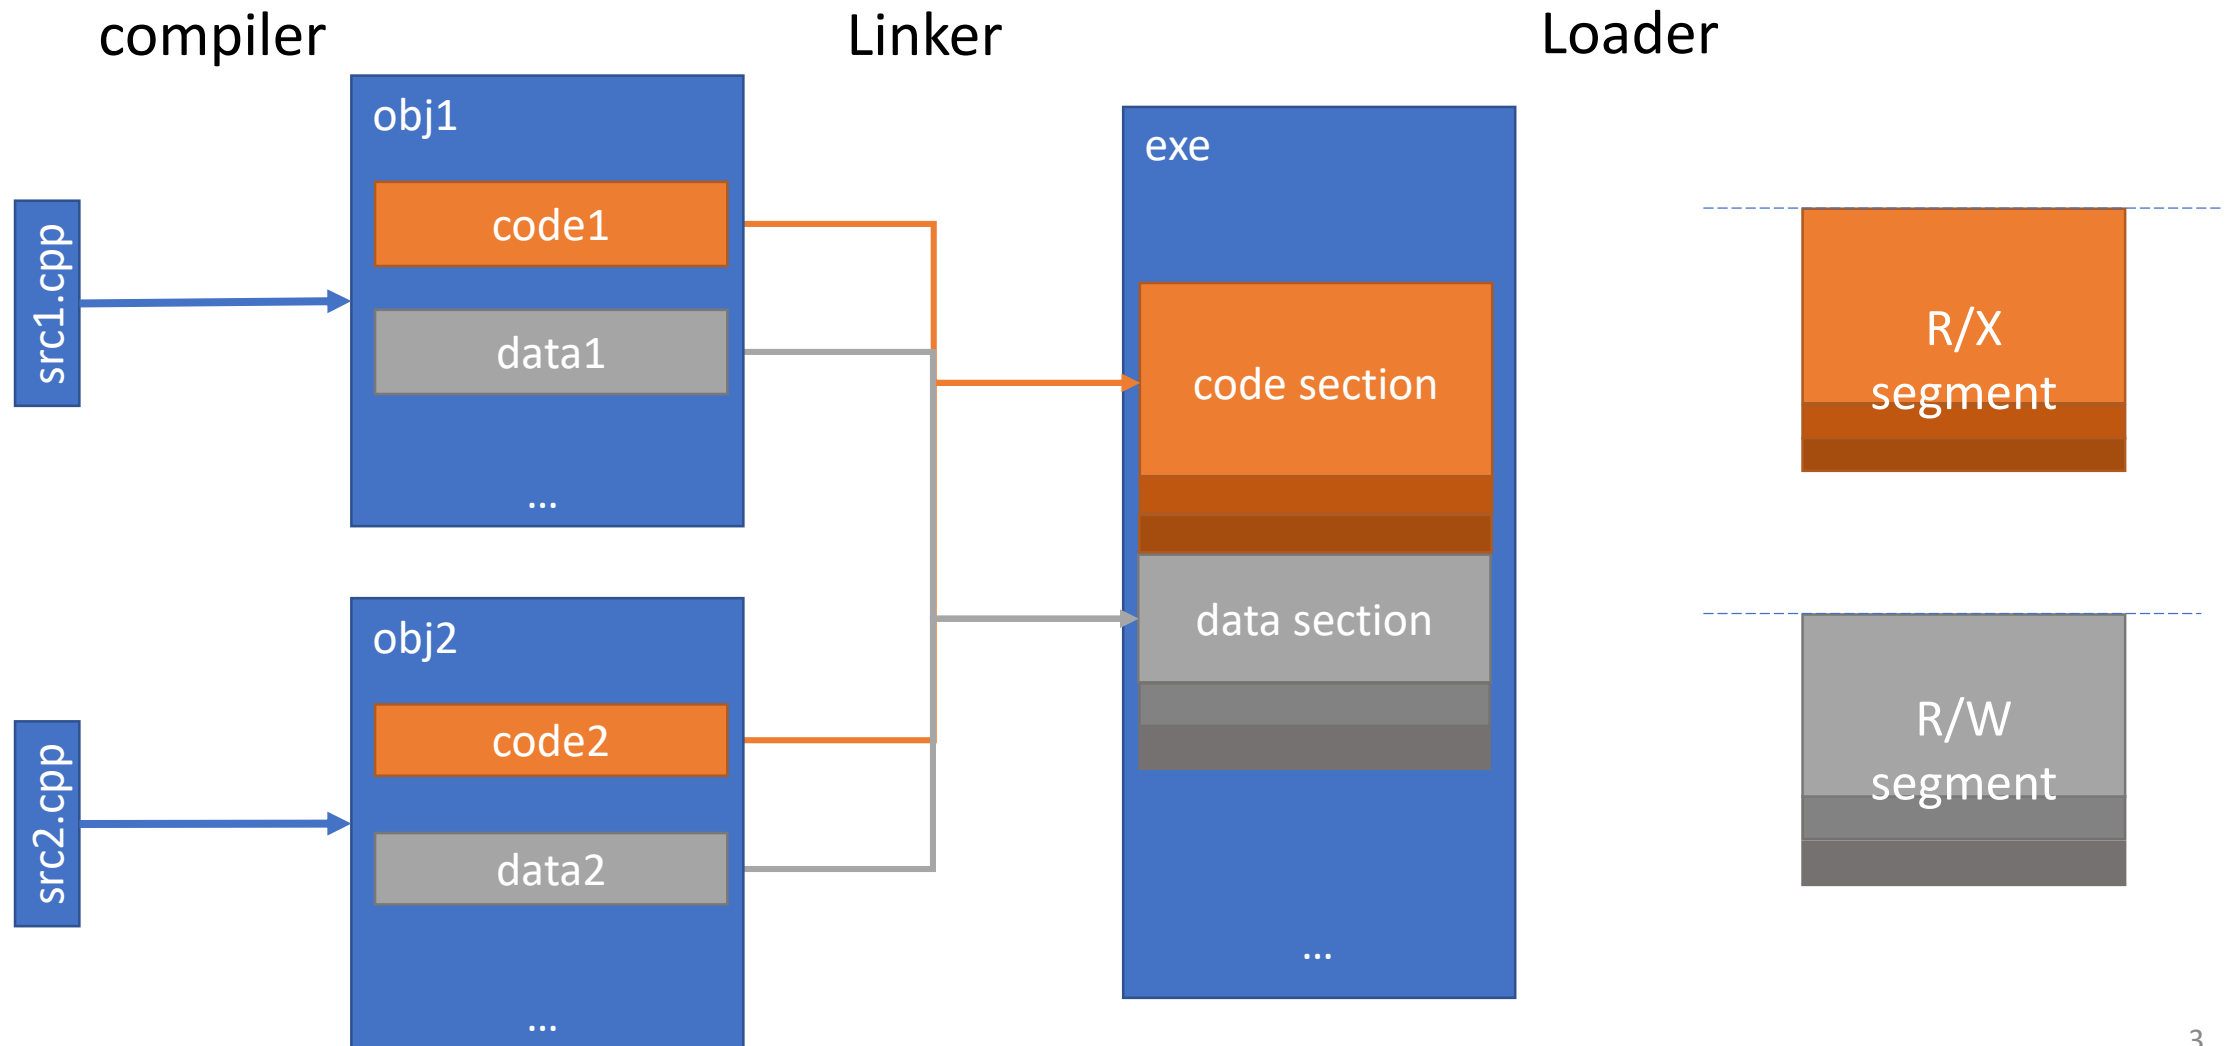
Senior Developer
@Toga Networks
(a Huawei Company)

ofekshilon@gmail.com




OfekShilon

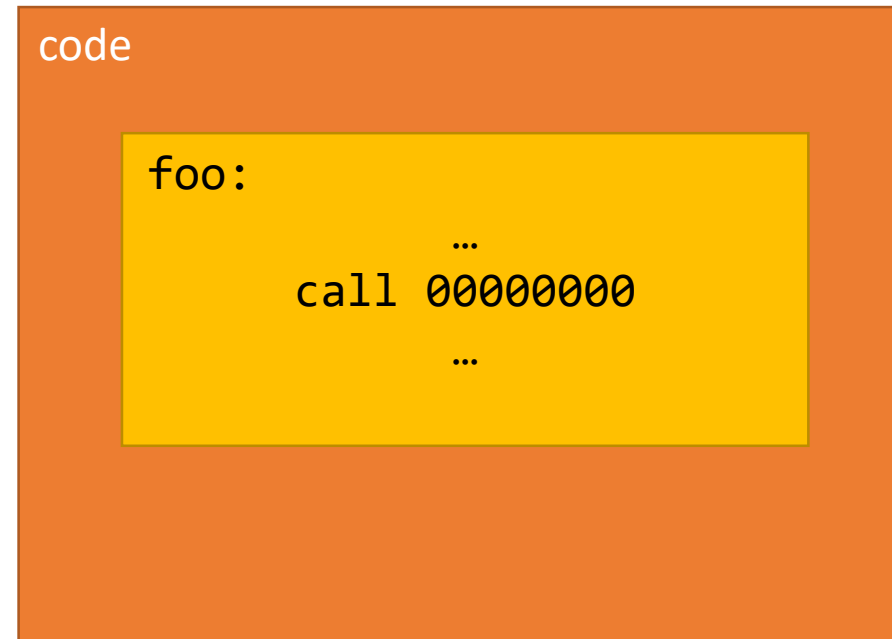
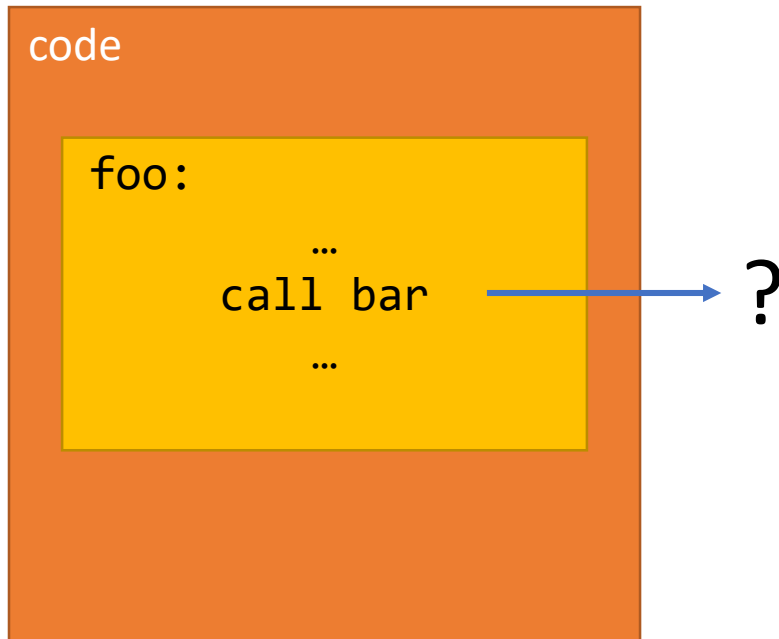
Intro to Linking in 3 slides, #1



Intro to Linking in 3 slides, #2



Intro to Linking in 3 slides, #3




.reloc


“Find ‘bar’ and write its address
over the 00000000 placeholder”

Small lie 1


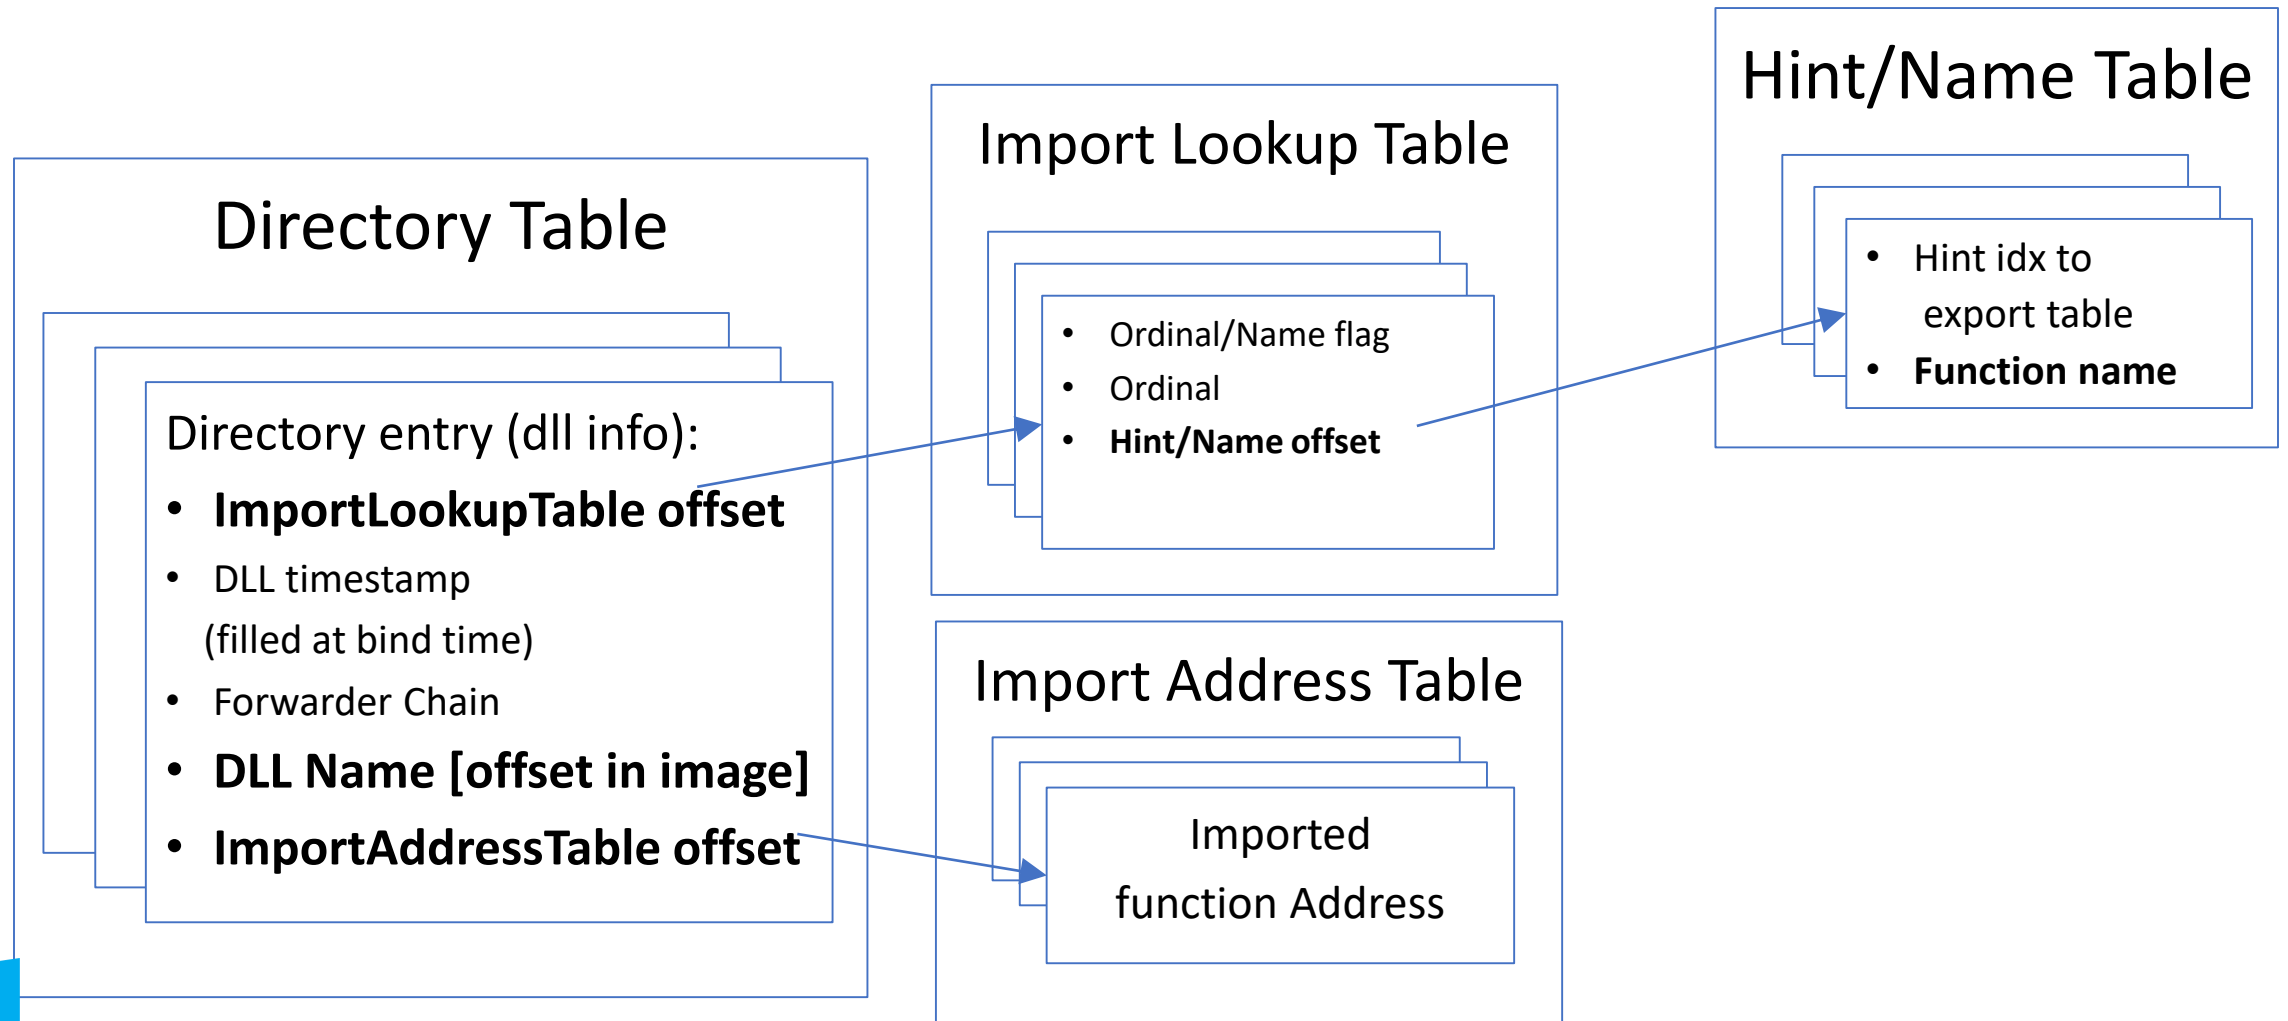
Windows



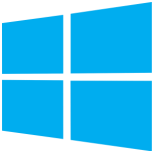
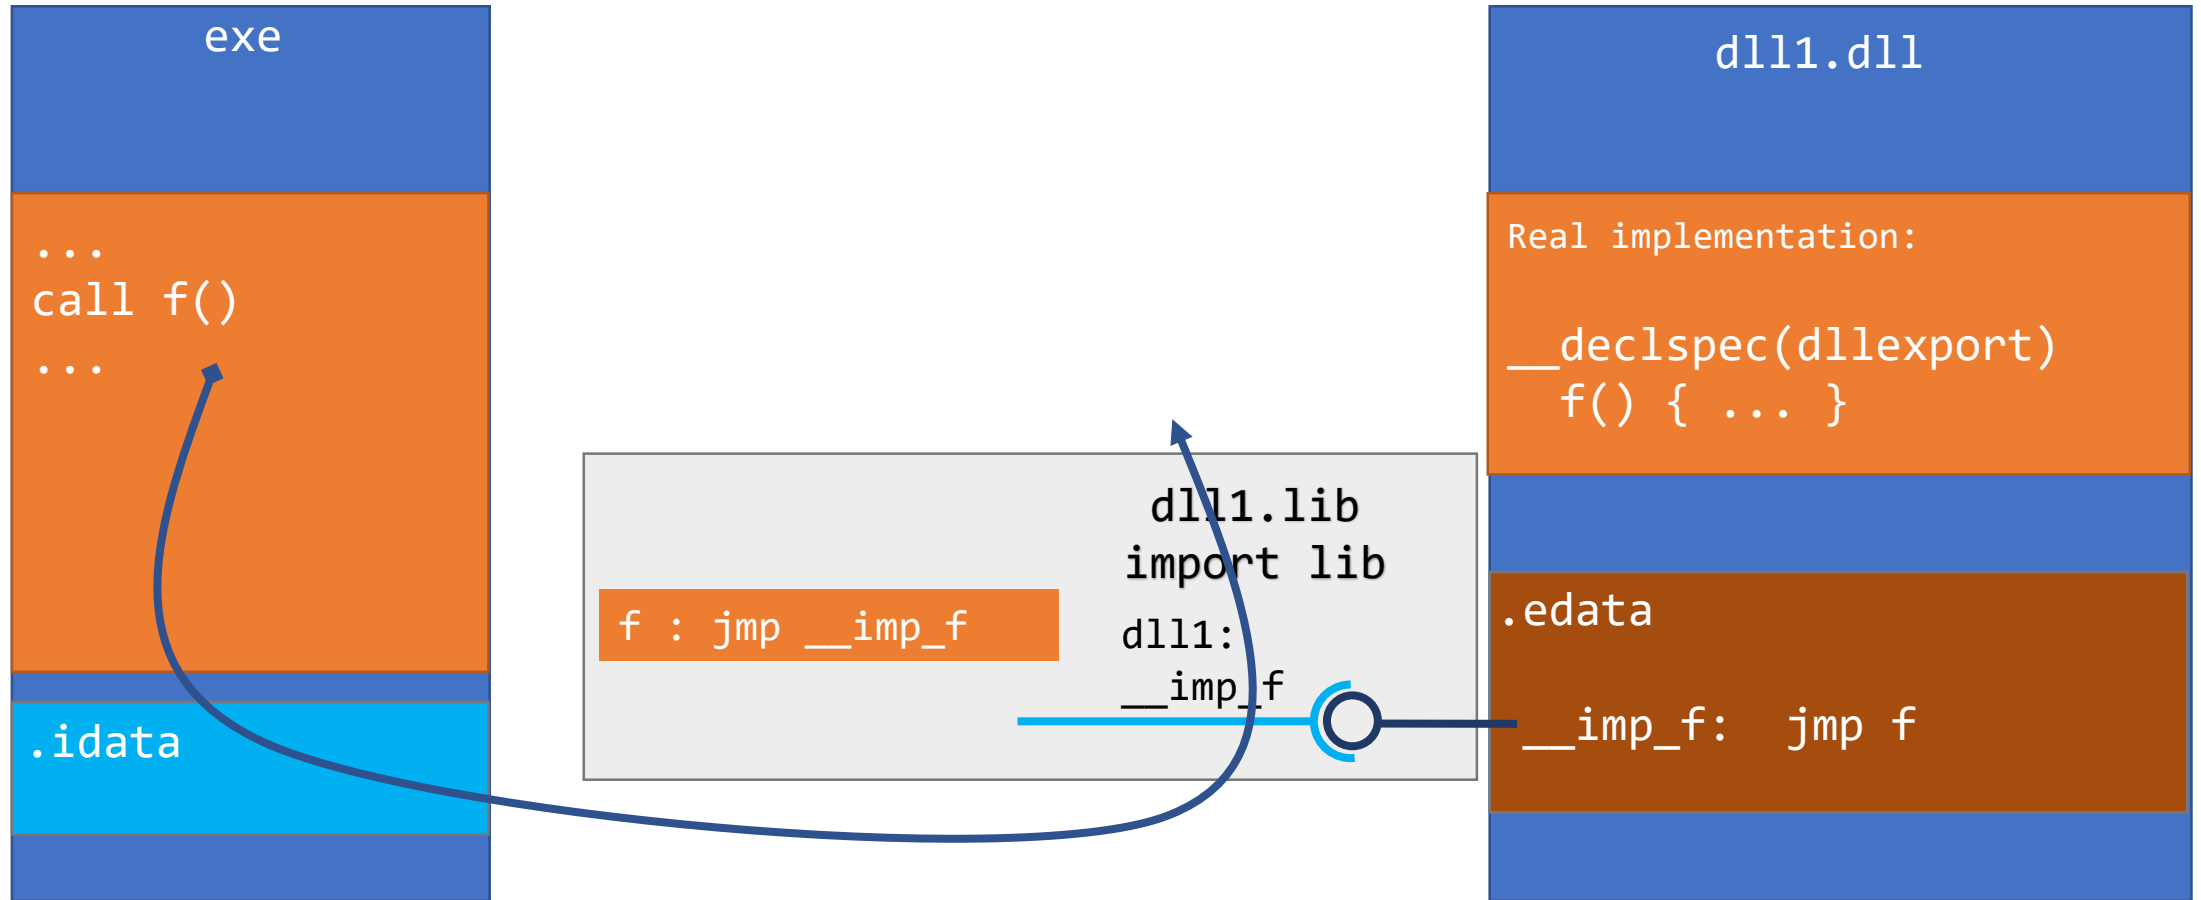
Import data section




.idata section layout



Import Library



Windows Schematic Interface



Linux



Linux import sections

- .dynamic /.dynsym: separate buckets of lib names and symbol names

```
Dynamic section at offset 0x21a58 contains 28 entries:
```


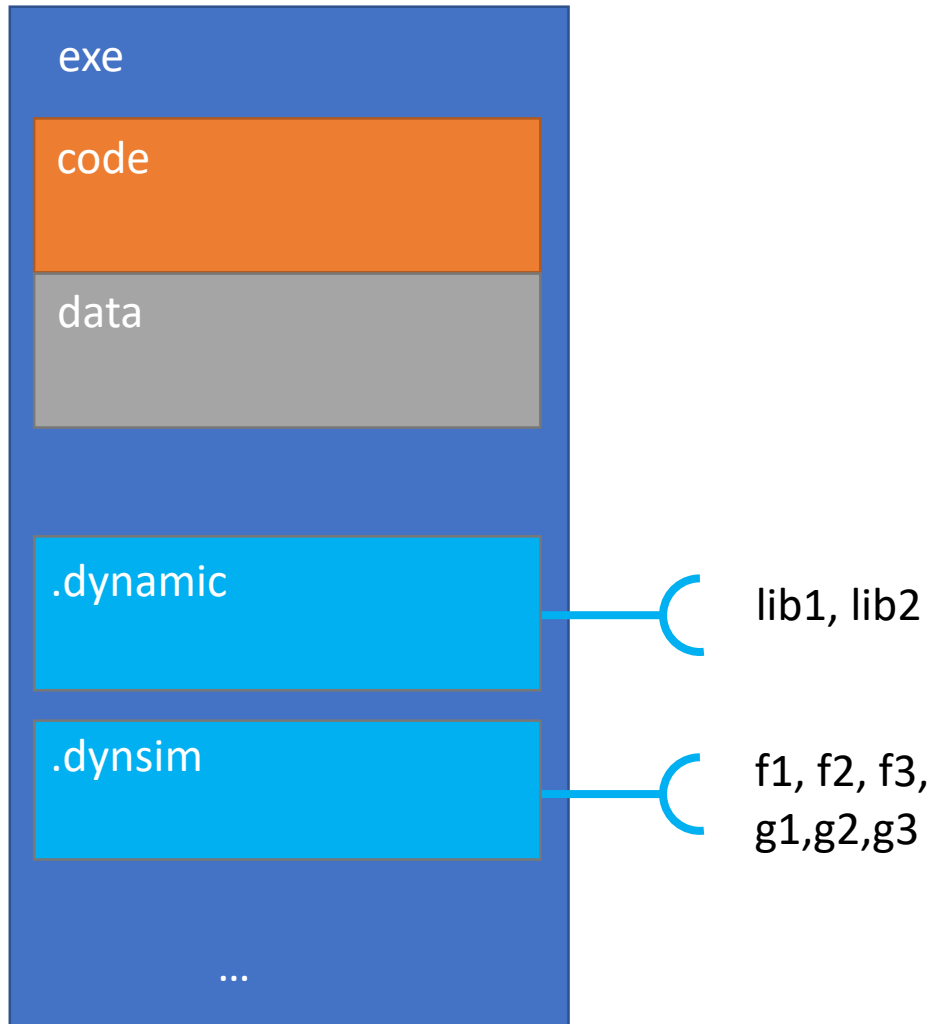
Tag	Type	Name/Value
0x0000000000000001	(NEEDED)	Shared library: [libselinux.so.1]
0x0000000000000001	(NEEDED)	Shared library: [libc.so.6]
0x000000000000000c	(INIT)	0x4000

```
Symbol table '.dynsym' contains 139 entries:
```

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
0:	0000000000000000	0	NOTYPE	LOCAL	DEFAULT	UND	
1:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	__ctype_toupper_loc@GLIBC_2.3 (2)
2:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	getenv@GLIBC_2.2.5 (3)
3:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	sigprocmask@GLIBC_2.2.5 (3)
4:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	__snprintf_chk@GLIBC_2.3.4 (4)



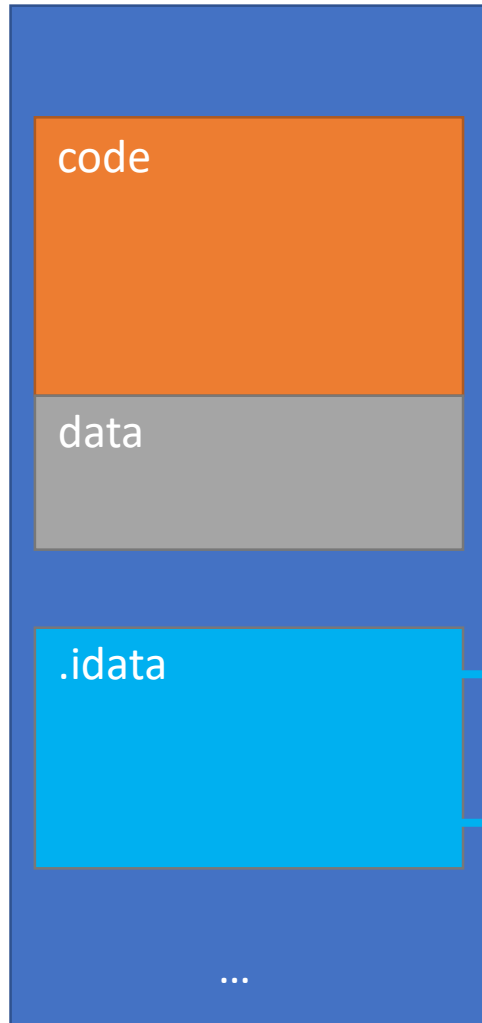
Linux Schematic Interface



Symbol search order:

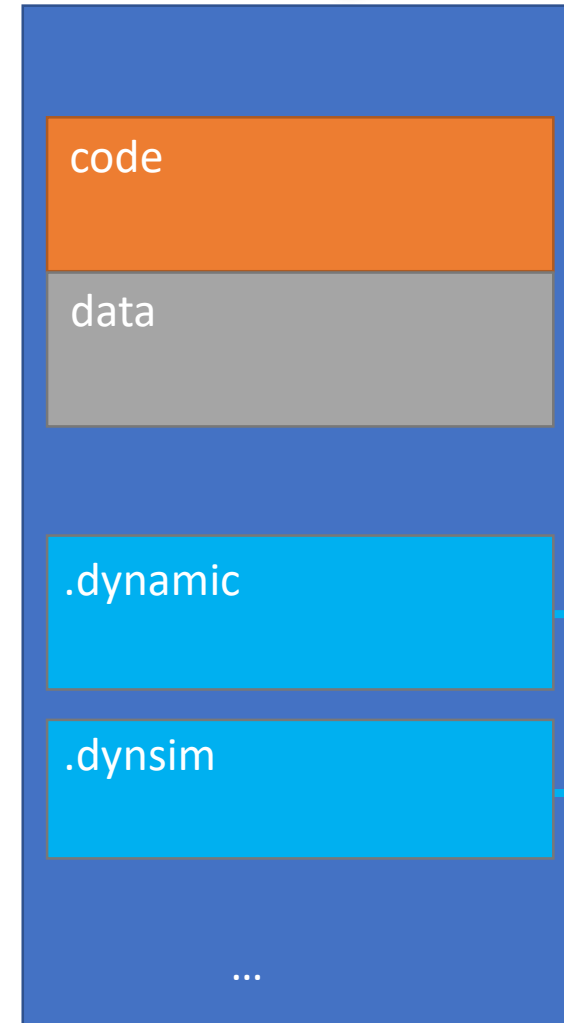
- By default, **exe before current lib**
- Controllable with -
 - **-Bsymbolic***,
 - **--dynamic-list***





lib1: f1, f2, f3

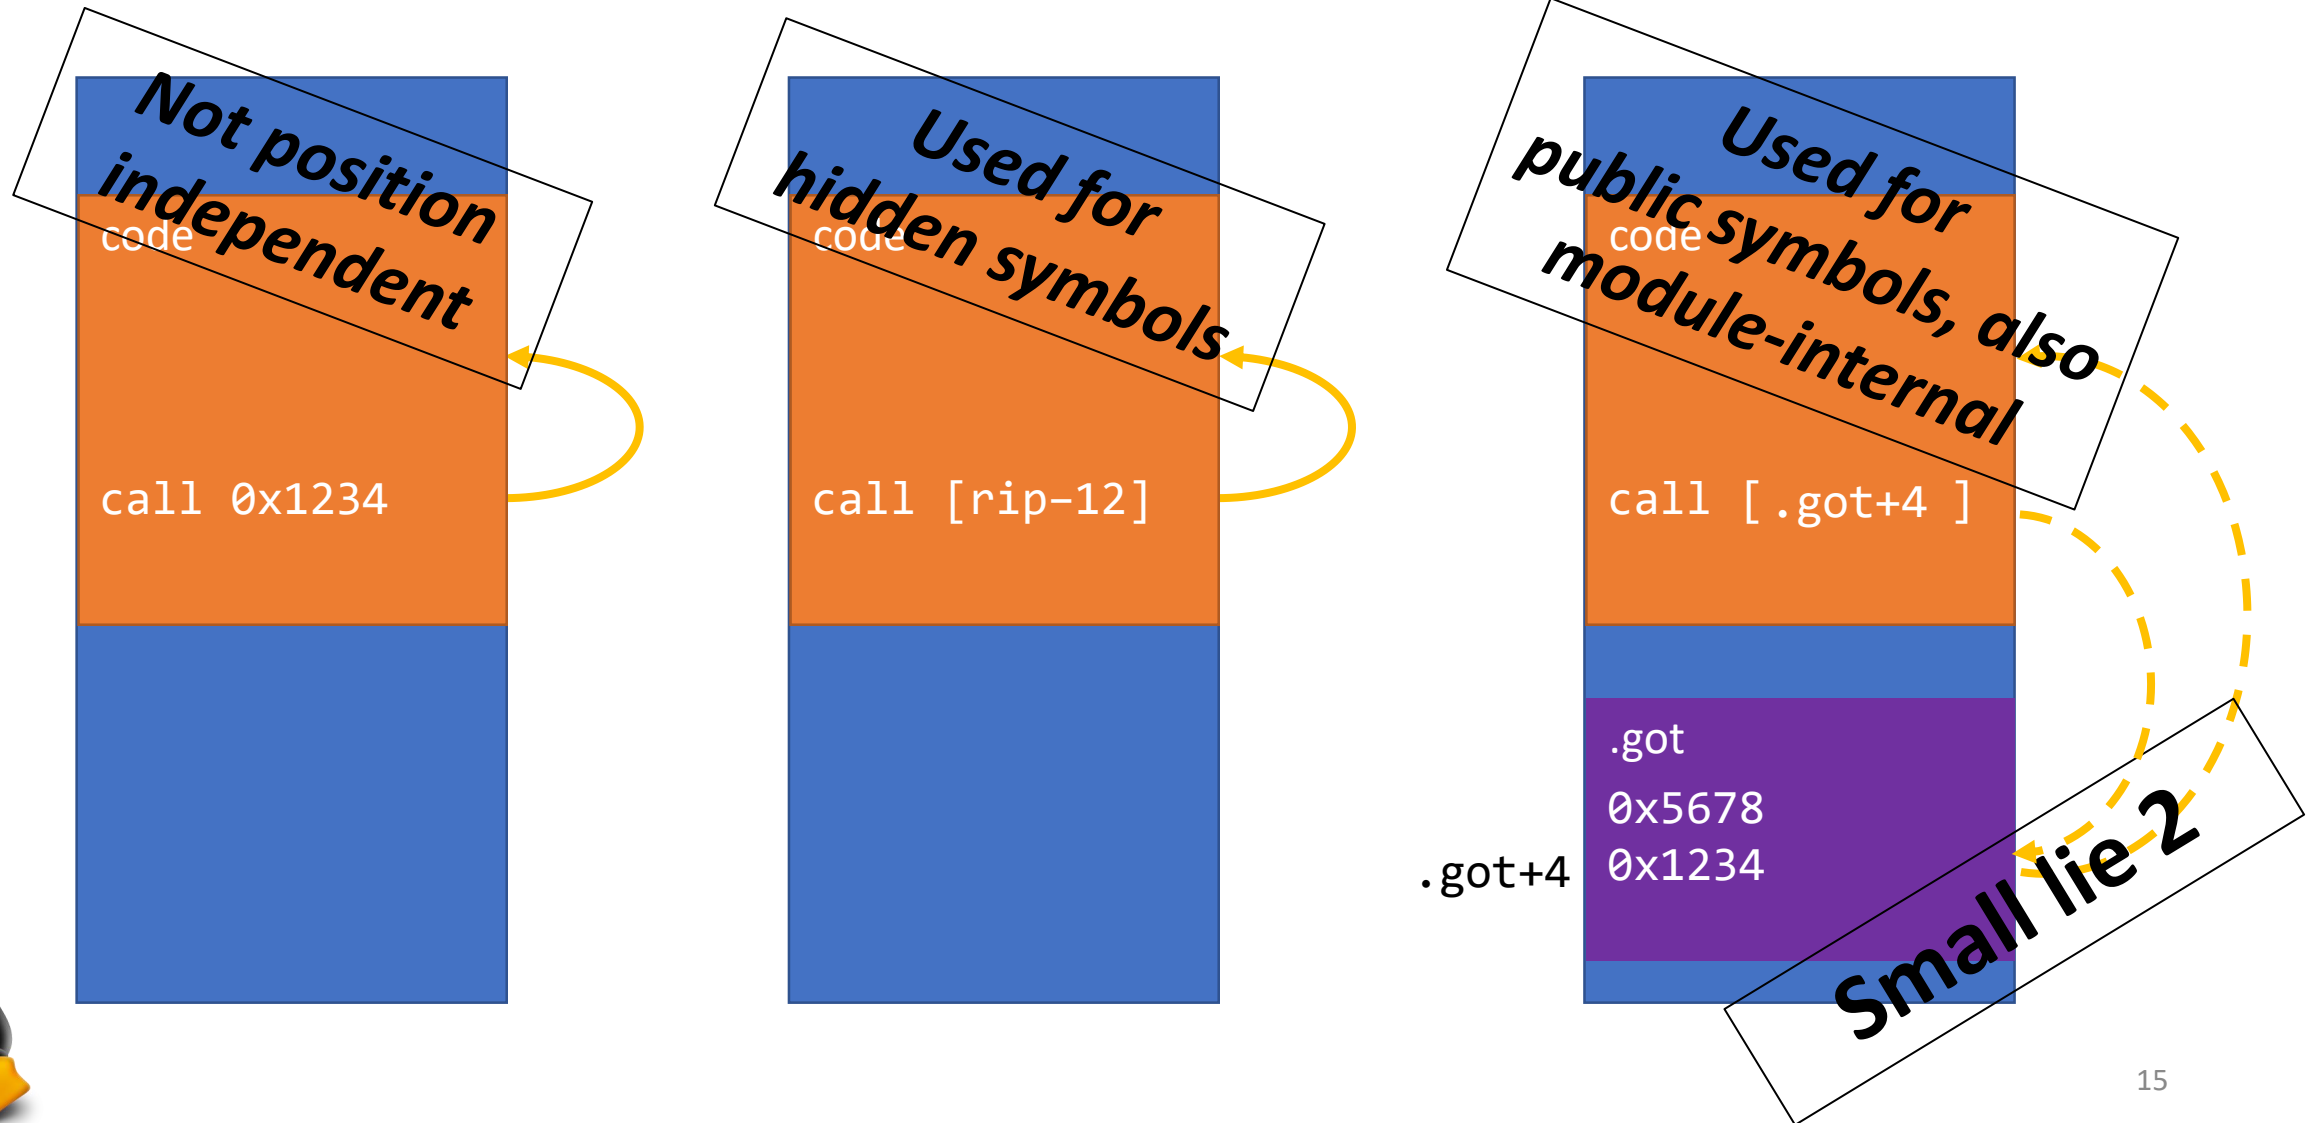
lib2: g1, g2, g3



lib1, lib2

f1, f2, f3,
g1,g2,g3

Position Independent Code

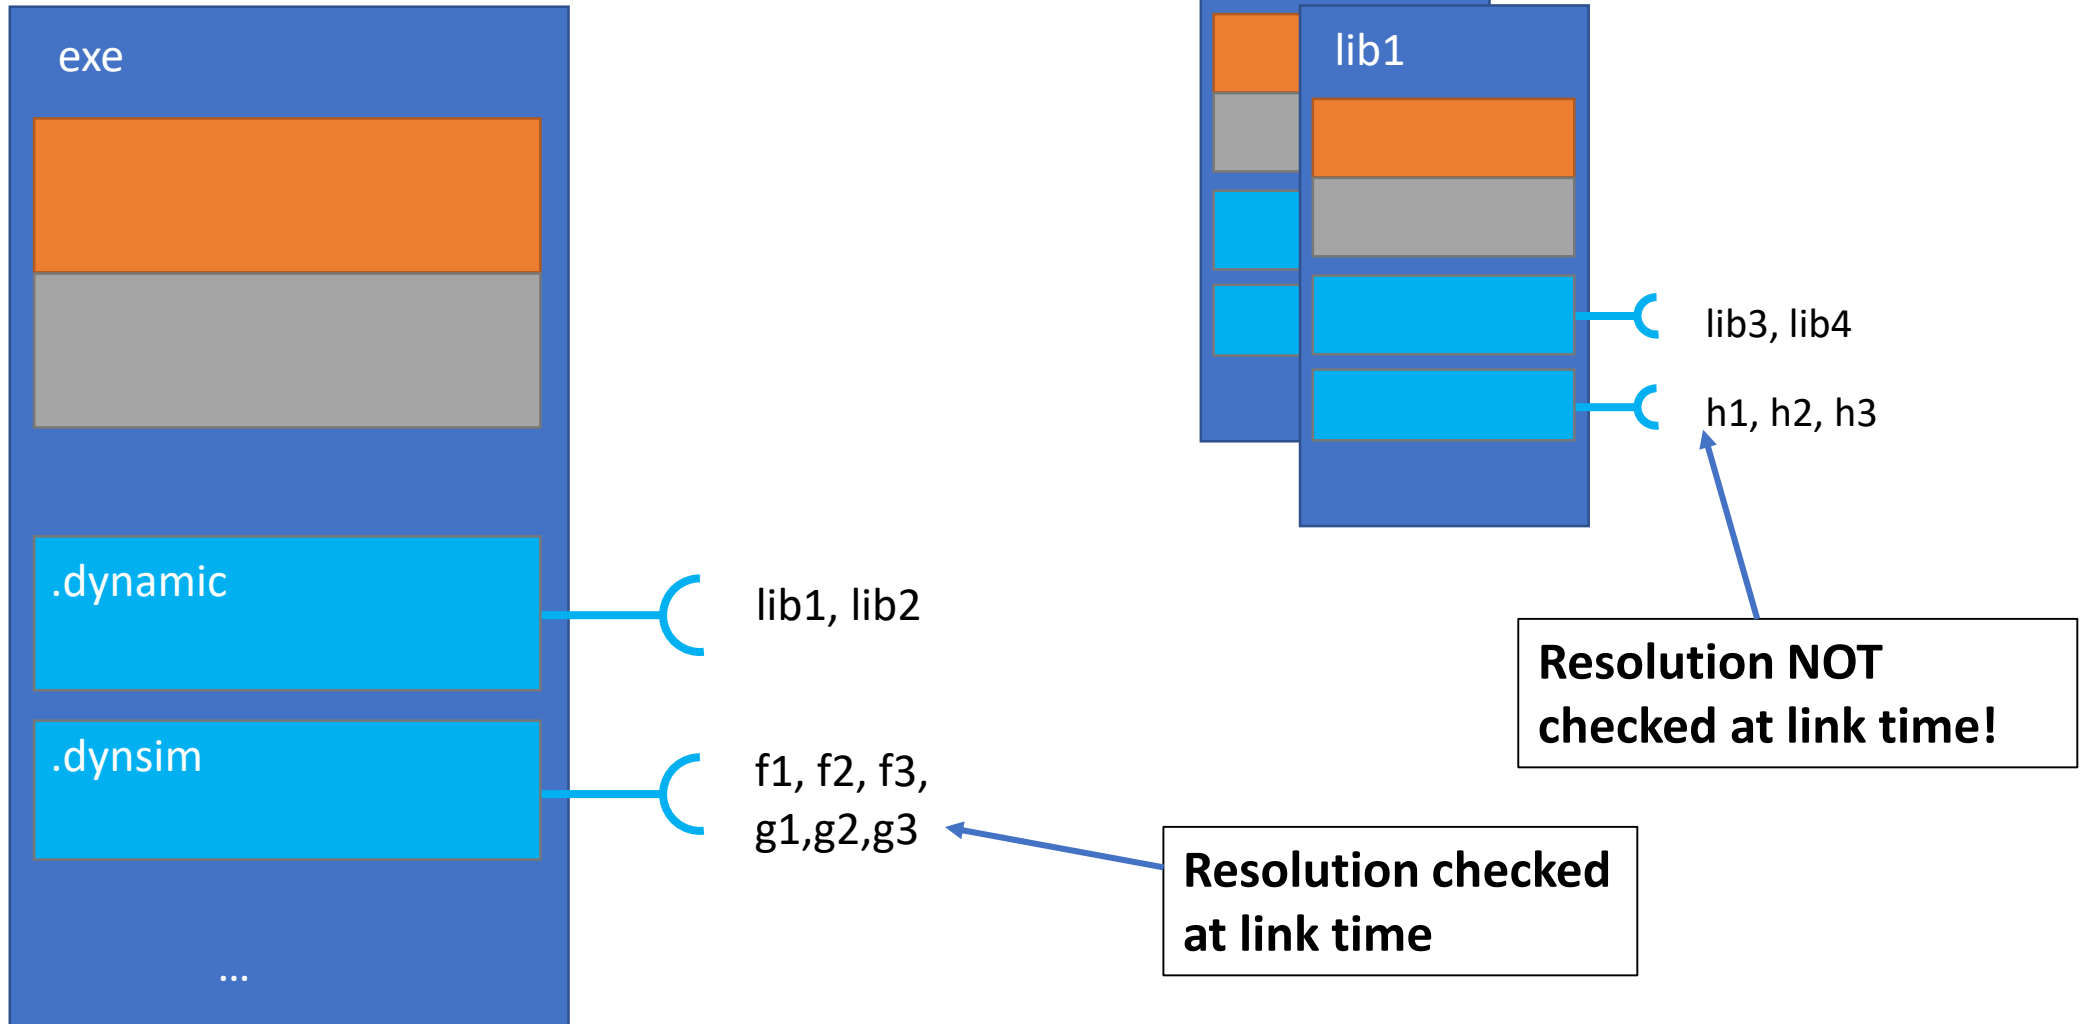


Position Independent Code - switches

- **ALL** shared lib code must be position independent
- Yet -fPIC is optional , not even implicit for -shared.
- If you try to link a shared lib from obj files not built with fPIC *and* using global vars:
`error: relocation R_X86_64_PC32 against symbol `global' can not be used when making a shared object; recompile with -fPIC`
- -fpic vs -fPIC:
 - Help bypass built-in GOT size limit on some architectures (alpha, sparc).
- -fPIC vs -fPIE:
 - <http://www.openbsd.org/papers/nycbsdcon08-pie/mgp00004.html>



Resolution Time



Resolution Time

- Default: `--allow-shlib-undefined`
- Can be controlled with –
 - `--no-allow-shlib-undefined`
 - Note: operates recursively on ld, not on gold / lld.
 - `-z defs /--no-undefined` : forces link time resolution check



Linux: Intermediate Summary

- Symbol and Library dependencies are maintained separately
- By default ***all*** calls are indirected through `.got`
- By default, resolution in libraries (essentially populating the `.got` in every library) is deferred to load time, **and** the executable is searched first.



C++ Implication #1: How to form a process-wide singleton?

- (variable or function)
- Linux:
 - Just put it in the executable
- Windows:
 - Re-link the EXE and all DLLs against the single DLL that defines the singleton.

C++ Implication #2: Can you have circular library dependencies?

- Linux:
 - Yes
- Windows
 - No.
 - Well, you'd have to hack hard.
- The Linux design provides some flexibility, but ..
 - “This [allowed-shlib-undefined] is an unfortunate default for -shared. Changing it may be disruptive today. Mach-O and PE/COFF have many problems but this may be a place where they got right.”
Rui Ueyama, author of LLD and MOLD linkers
<https://maskray.me/blog/2021-06-13-dependency-related-linker-options>

C++ Implication #3: Can a shared-library symbol be overridden from an executable?

- “Interposition”
- Windows:
 - No.
 - Well –
 - Not from the executable. You’d have to re-link all components against a dll implementing the symbol to be overridden.
- Linux:
 - Yes.
 - Non-default build switches can intervene.

C++: new

[replacement.functions]: A C++ program may provide the definition for any of the following dynamic memory allocation function signatures declared in header `<new>` :

- `operator new(std::size_t)`
- `operator new(std::size_t, std::align_val_t)`
- ...

The program's definitions are used instead of the default versions supplied by the implementation ...

- On Windows, that's not what happens.

Lazy Binding


(a.k.a Delayed Loading)

Lazy Bind by Default?



Yes.

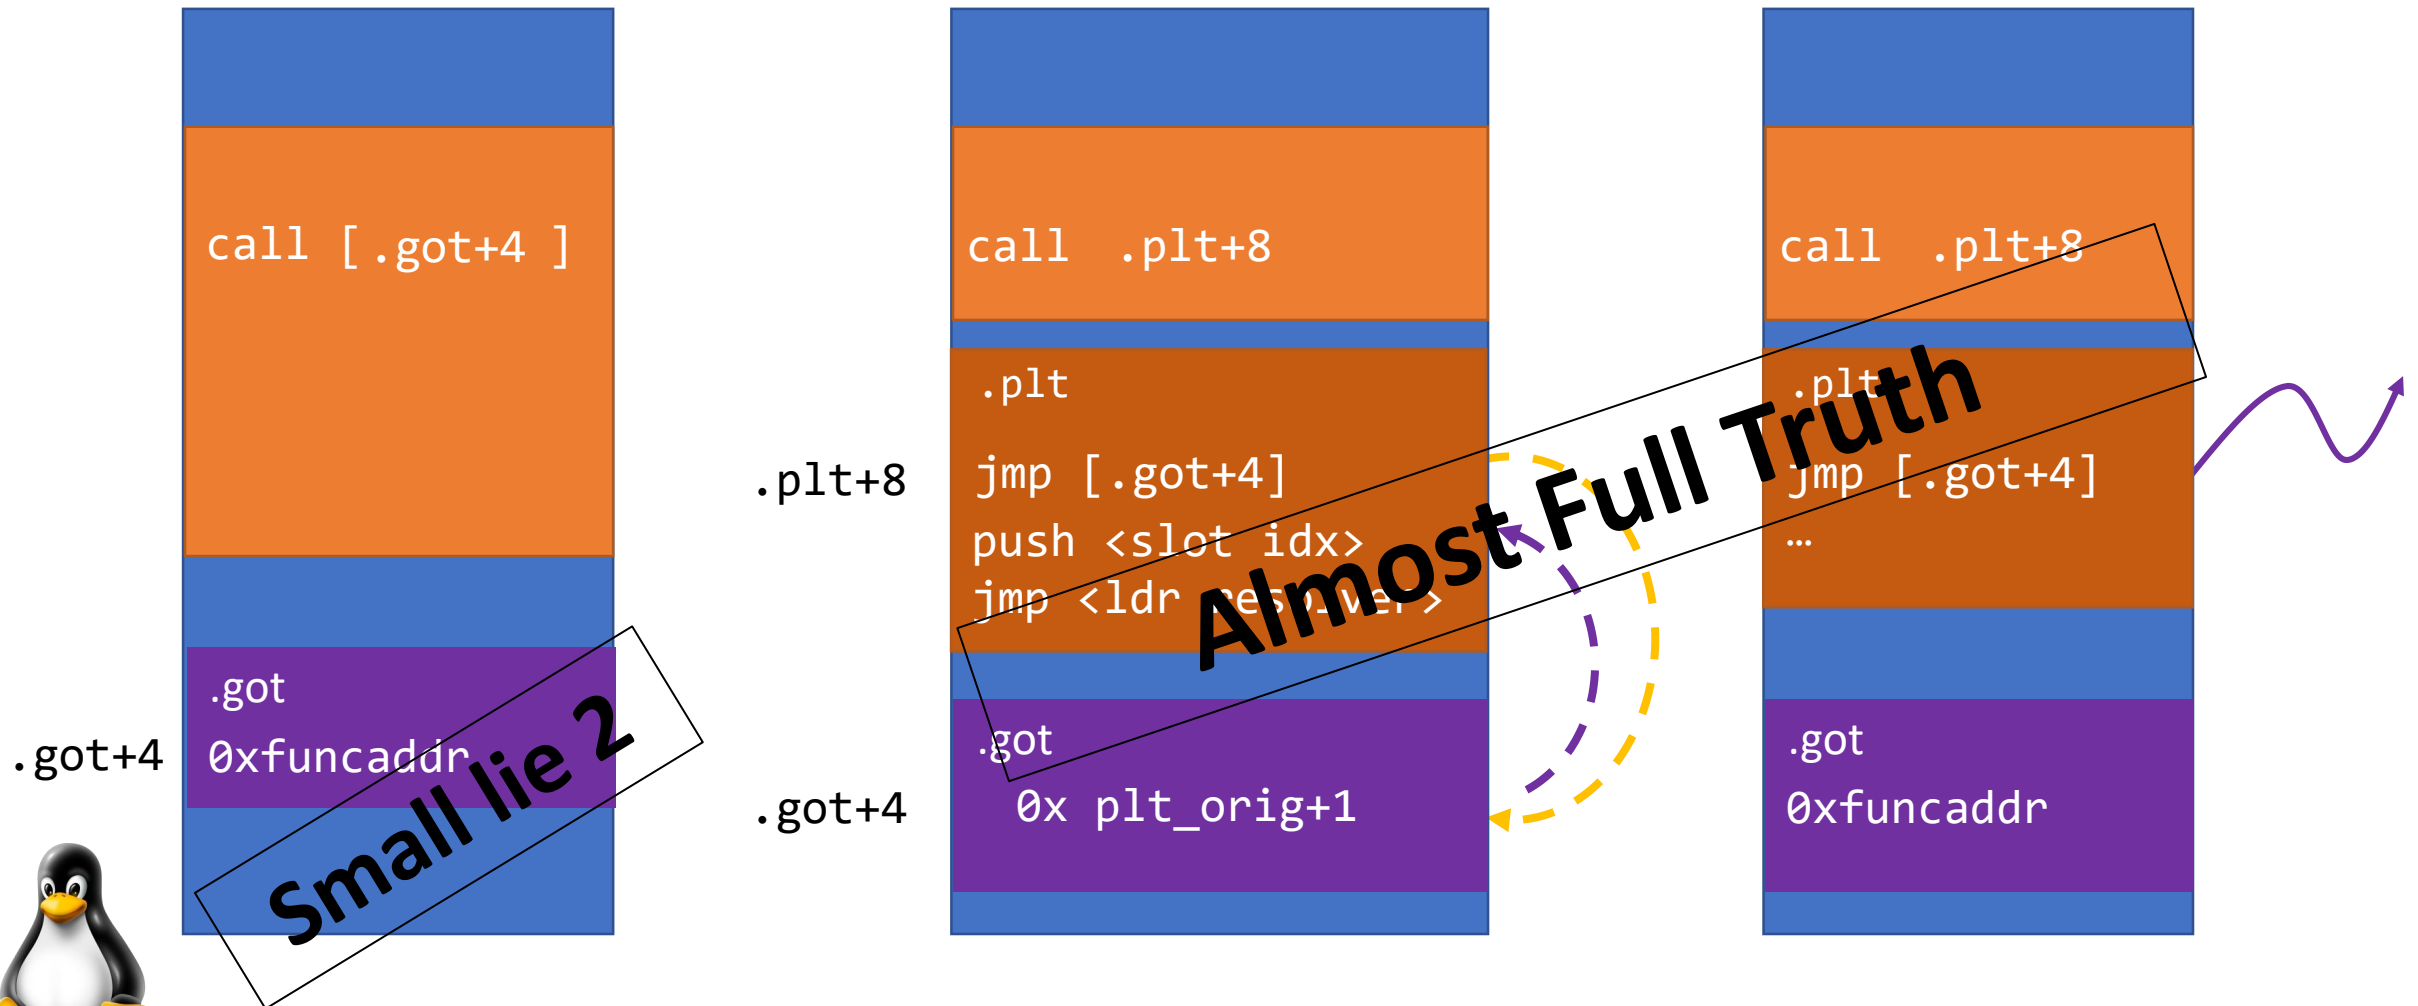
Controllable with the env var
`LD_BIND_NOW`



No.

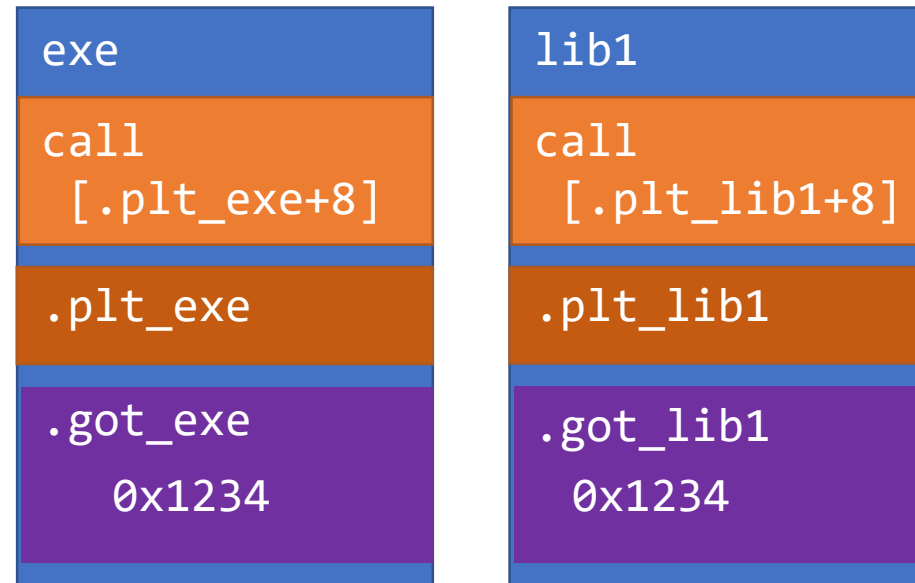
Controllable with linker switch
`/DELAYLOAD:<your_dll.dll>`

Procedure Linkage Table (PLT)



C++ Implication #4: Comparing Func Ptrs

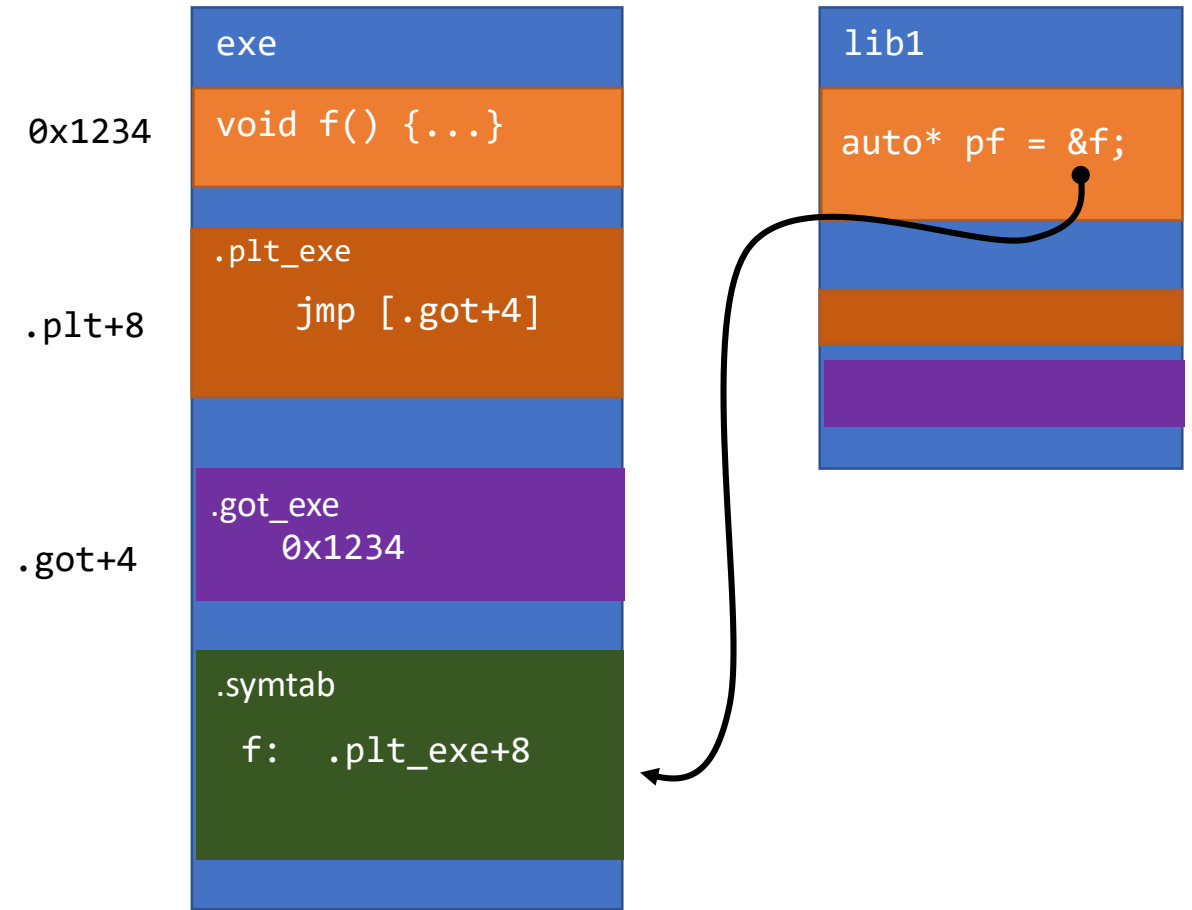
- C++ standard, [expr.eq]§3.2: “... if the pointers are both null, **both point to the same function**, or both represent the same address (6.8.2), they compare equal.”
- Actual calls are made to a PLT entry.
- Different among libraries!



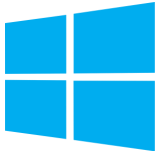
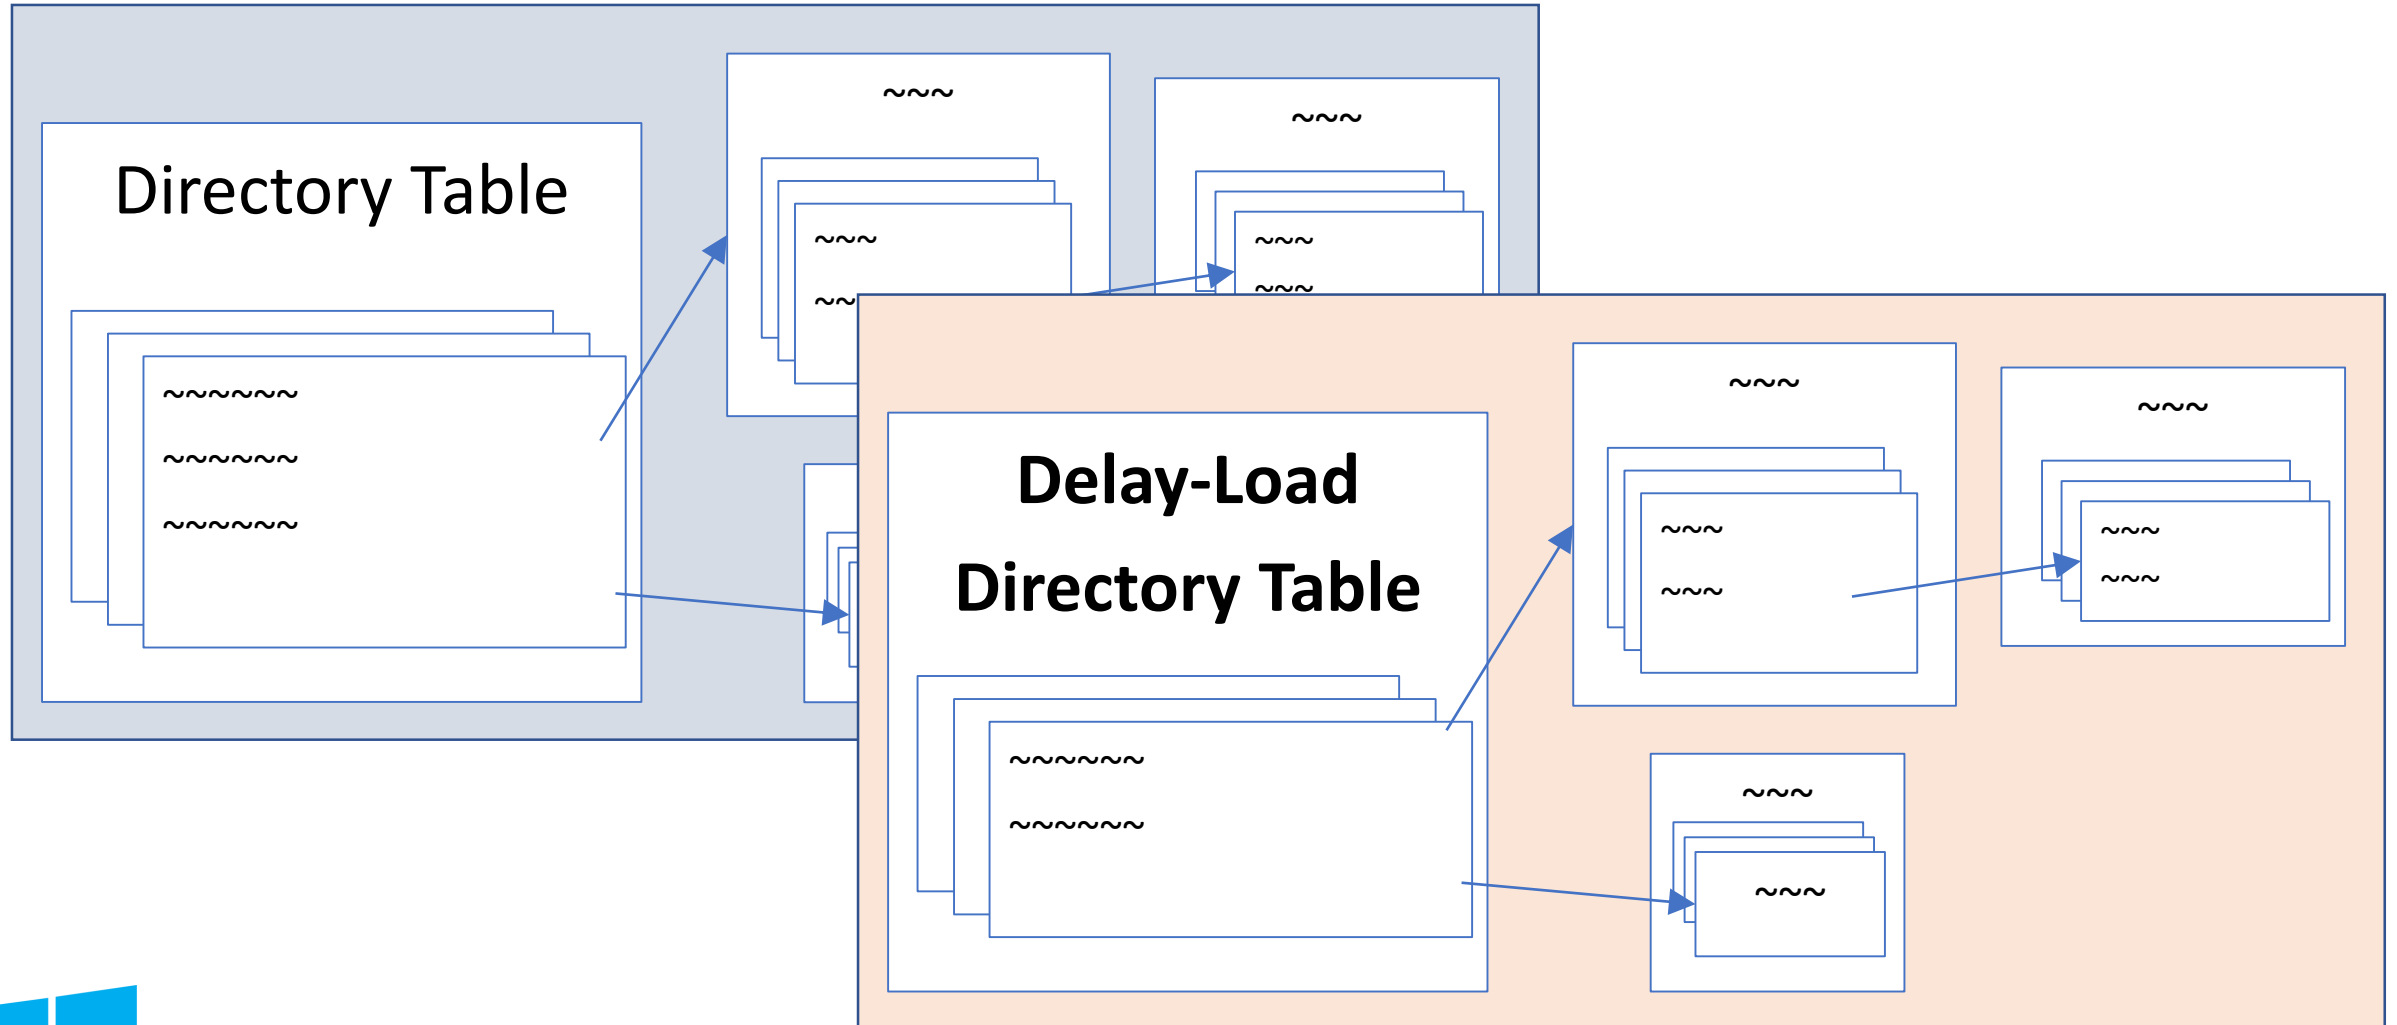
C++ Implication #4: Comparing Func Ptrs

From the SystemV ABI:

To allow comparisons of function addresses to work as expected, if an executable file references a function defined in a shared object, the link editor will place the address of the procedure linkage table entry for that function in its associated symbol table entry. This will result in symbol table entries with section index of SHN_UNDEF but a type of STT_FUNC and a non-zero st_value. A reference to the address of a function from within a shared library will be satisfied by such a definition in the executable



Windows .idata section



No analogue mechanism in Windows:

```
1
2  #include <iostream>
3  using std::cout;
4
5  __declspec(dllimport) void delayedLoadFunc();
6
7  int main()
8  {
9      cout << &delayedLoadFunc << "\n";
10     delayedLoadFunc();
11     cout << &delayedLoadFunc << "\n";
12 }
13
14
```

Microsoft Visual Studio Debug Console

00007FF79875149C
00007FFBB25211D1


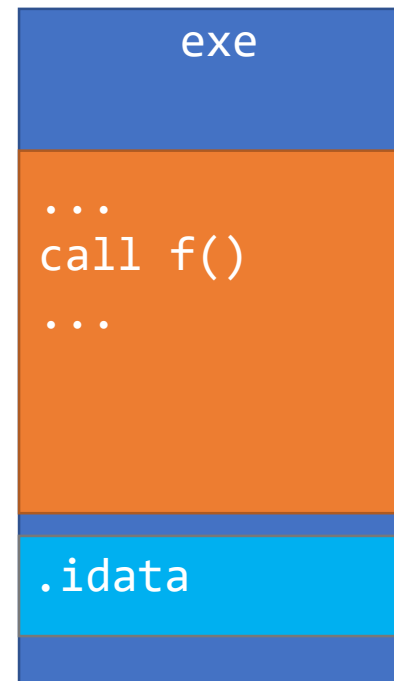
C:\Users\o00836696\source\repos\testDelayLoad
Press any key to close this window . . .



Symbol Visibility

Symbol Visibility - Windows

- `__declspec(dllexport)` – add symbol to .edata
- `__declspec(import)` – doesn't do much...
 - Minor optimization that happens in Release anyway
- Most symbols are neither.



Symbol Visibility - Linux

- No import/export distinction! Only public/hidden
- When a symbol is public it –
 - Goes through the longer PLT/GOT route,
 - Is available to other libraries/executables (“exported”).
 - Potentially subject to interposition (“imported”), and allowed-undefined.
- ***Default visibility is public!***
 - Intervene through:
 - -fvisibility=hidden,
 - -fvisibility-inlines-hidden,
 - -fvisibility-ms-compat,
 - `__attribute__((visibility("hidden")))`



Symbol Visibility - Linux

- “Using this feature [-fvisibility=hidden] can very substantially improve linking and load times of shared object libraries, produce more optimized code, provide near-perfect API export and prevent symbol clashes. It is strongly recommended that you use this in any shared objects you distribute.”

\$ man gcc

- <https://gcc.gnu.org/wiki/Visibility>



Virtual functions overhead

- Careful when measuring on linux...
- By default **ALL** calls are indirect calls, through the PLT/GOT.
 - In a library.
 - Sometimes in executables too (-fPIE)

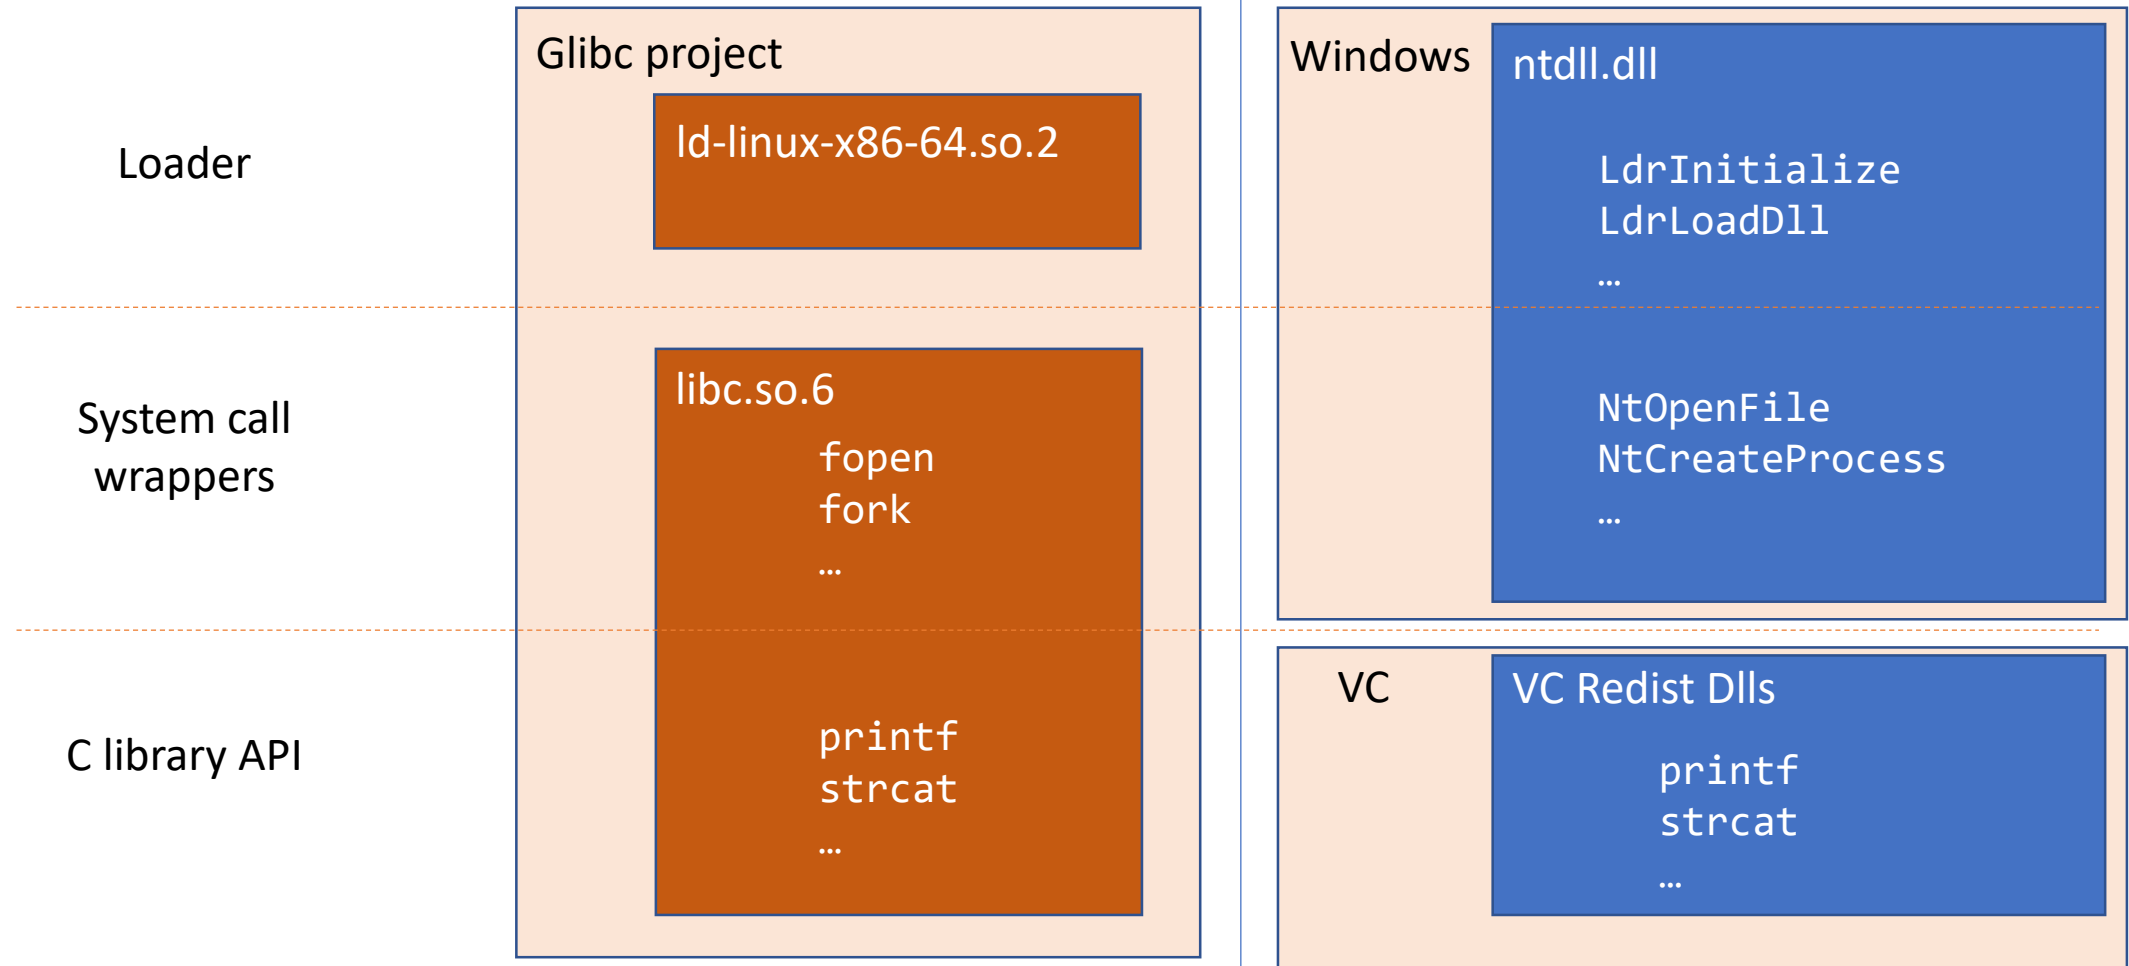


Loader

Loader

- A.k.a Dynamic Linker, a.k.a Interpreter, a.k.a Image Loader (esp in windows)
- Runs in user mode – operates on regular process address space

Component Map



Selecting a different loader

```
$ gcc -v whatever.cpp
...
COLLECT_GCC_OPTIONS='-v' '-mtune=generic' '-march=x86-64'
 /usr/lib/gcc/x86_64-linux-gnu/9/collect2 -plugin /usr/lib/gcc/x86_64-linux-gnu/9/liblto_plugin.so -
plugin-opt=/usr/lib/gcc/x86_64-linux-gnu/9/lto-wrapper -plugin-opt=-fresolution=/tmp/ccfheLJQ.res -
plugin-opt=-pass-through=-lgcc -plugin-opt=-pass-through=-lgcc_s -plugin-opt=-pass-through=-lc -
plugin-opt=-pass-through=-lgcc -plugin-opt=-pass-through=-lgcc_s --build-id --eh-frame-hdr -m
elf_x86_64 --hash-style=gnu --as-needed -dynamic-linker /lib64/ld-linux-x86-64.so.2 -pie -z now -z
relro /usr/lib/gcc/x86_64-linux-gnu/9/../../../../x86_64-linux-gnu/Scrt1.o /usr/lib/gcc/x86_64-linux-
gnu/9/../../../../x86_64-linux-gnu/crti.o /usr/lib/gcc/x86_64-linux-gnu/9/crtbeginS.o -
L/usr/lib/gcc/x86_64-linux-gnu/9 -L/usr/lib/gcc/x86_64-linux-gnu/9/../../../../x86_64-linux-gnu -
L/usr/lib/gcc/x86_64-linux-gnu/9/../../../../lib -L/lib/x86_64-linux-gnu -L/lib/../../lib -
L/usr/lib/x86_64-linux-gnu -L/usr/lib/../../lib -L/usr/lib/gcc/x86_64-linux-gnu/9/../../../../
/tmp/ccijfiIQ.o -lgcc --push-state --as-needed -lgcc_s --pop-state -lc -lgcc --push-state --as-
needed -lgcc_s --pop-state /usr/lib/gcc/x86_64-linux-gnu/9/crtendS.o /usr/lib/gcc/x86_64-linux-
gnu/9/../../../../x86_64-linux-gnu/crtn.o
...
```



Selecting a different loader

```
$ readelf --program-headers /usr/bin/ls
```

...

Program Headers:

Type	Offset FileSiz	VirtAddr MemSiz	PhysAddr Flags	Align
PHDR	0x0000000000000040	0x0000000000000040	0x0000000000000040	
	0x00000000000002d8	0x00000000000002d8	R	0x8
INTERP	0x0000000000000318	0x0000000000000318	0x0000000000000318	
	0x000000000000001c	0x000000000000001c	R	0x1

[Requesting program interpreter: /lib64/ld-linux-x86-64.so.2]

...



Observing the Loader in action - Linux

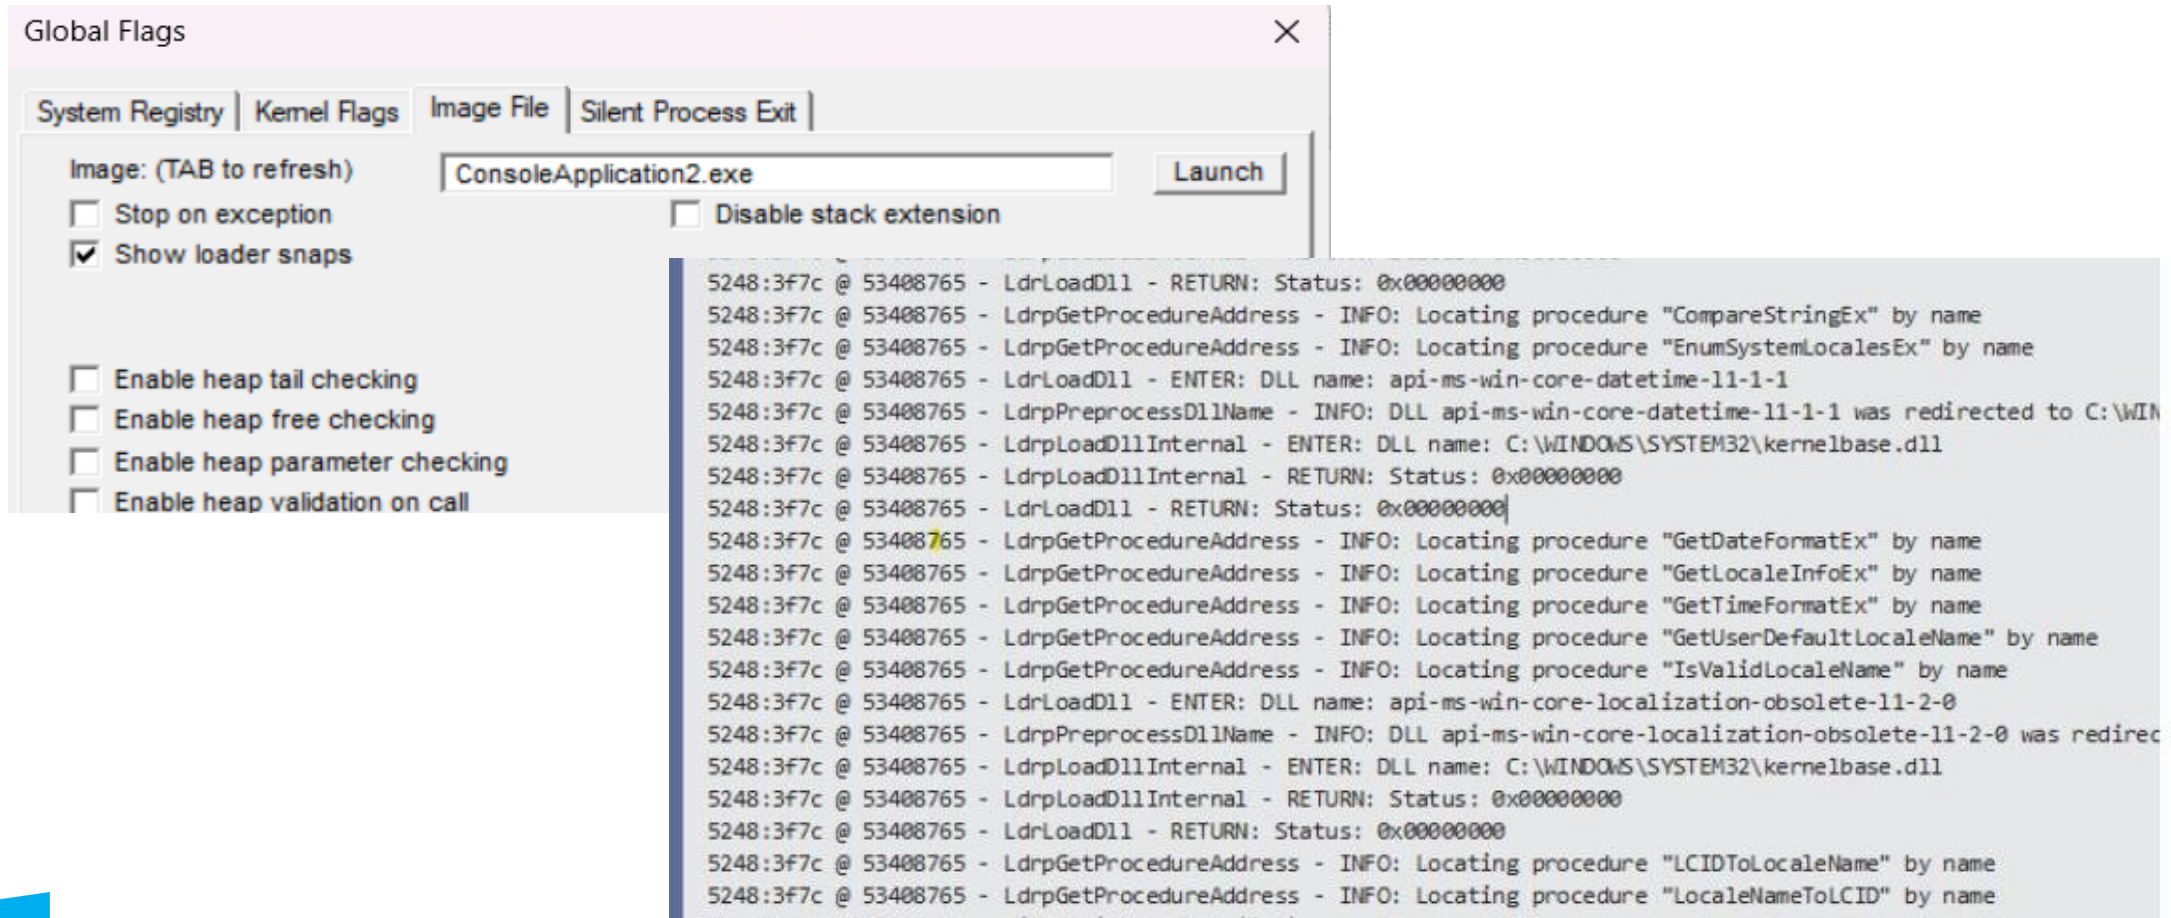
```
$ LD_DEBUG=help cat
```

Valid options for the LD_DEBUG environment variable are:

libs	display library search paths
reloc	display relocation processing
files	display progress for input file
symbols	display symbol table processing
bindings	display information about symbol binding
versions	display version dependencies
scopes	display scope information
all	all previous options combined
statistics	display relocation statistics
unused	determined unused DSOs
help	display this help message and exit



Observing the Loader in action - Windows



The screenshot displays the 'Global Flags' dialog box with the 'Image File' tab selected. The 'Image' field contains 'ConsoleApplication2.exe' and the 'Launch' button is visible. Under the 'Image File' tab, the 'Show loader snaps' checkbox is checked, while 'Stop on exception' and 'Disable stack extension' are unchecked. Below this, there are five unchecked checkboxes: 'Enable heap tail checking', 'Enable heap free checking', 'Enable heap parameter checking', and 'Enable heap validation on call'.

Below the dialog box, a log window shows the following output:

```
5248:3f7c @ 53408765 - LdrLoadDll - RETURN: Status: 0x00000000
5248:3f7c @ 53408765 - LdrpGetProcedureAddress - INFO: Locating procedure "CompareStringEx" by name
5248:3f7c @ 53408765 - LdrpGetProcedureAddress - INFO: Locating procedure "EnumSystemLocalesEx" by name
5248:3f7c @ 53408765 - LdrLoadDll - ENTER: DLL name: api-ms-win-core-datetime-l1-1-1
5248:3f7c @ 53408765 - LdrpPreprocessDllName - INFO: DLL api-ms-win-core-datetime-l1-1-1 was redirected to C:\WIN
5248:3f7c @ 53408765 - LdrpLoadDllInternal - ENTER: DLL name: C:\WINDOWS\SYSTEM32\kernelbase.dll
5248:3f7c @ 53408765 - LdrpLoadDllInternal - RETURN: Status: 0x00000000
5248:3f7c @ 53408765 - LdrLoadDll - RETURN: Status: 0x00000000
5248:3f7c @ 53408765 - LdrpGetProcedureAddress - INFO: Locating procedure "GetDateFormatEx" by name
5248:3f7c @ 53408765 - LdrpGetProcedureAddress - INFO: Locating procedure "GetLocaleInfoEx" by name
5248:3f7c @ 53408765 - LdrpGetProcedureAddress - INFO: Locating procedure "GetTimeFormatEx" by name
5248:3f7c @ 53408765 - LdrpGetProcedureAddress - INFO: Locating procedure "GetUserDefaultLocaleName" by name
5248:3f7c @ 53408765 - LdrpGetProcedureAddress - INFO: Locating procedure "IsValidLocaleName" by name
5248:3f7c @ 53408765 - LdrLoadDll - ENTER: DLL name: api-ms-win-core-localization-obsolete-l1-2-0
5248:3f7c @ 53408765 - LdrpPreprocessDllName - INFO: DLL api-ms-win-core-localization-obsolete-l1-2-0 was redirec
5248:3f7c @ 53408765 - LdrpLoadDllInternal - ENTER: DLL name: C:\WINDOWS\SYSTEM32\kernelbase.dll
5248:3f7c @ 53408765 - LdrpLoadDllInternal - RETURN: Status: 0x00000000
5248:3f7c @ 53408765 - LdrLoadDll - RETURN: Status: 0x00000000
5248:3f7c @ 53408765 - LdrpGetProcedureAddress - INFO: Locating procedure "LCIDToLocaleName" by name
5248:3f7c @ 53408765 - LdrpGetProcedureAddress - INFO: Locating procedure "LocaleNameToLCID" by name
```



C++ and Shared Libs

C++ and Shared Libs

[replacement.functions]: A C++ program may provide the definition for any of the following dynamic memory allocation function signatures declared in header `<new>` :

- `operator new(std::size_t)`
- `operator new(std::size_t, std::align_val_t)`
- ...

The program's definitions are used instead of the default versions supplied by the implementation

- “Shared libraries are out of scope for the standard”
- “ ‘Program’ and ‘implementation’ are undefined”

C++ and Shared Libs

- These clauses are dead letter.
 - Pragmatic suggestion: drop altogether statements that will never be followed.
 - replacement-function
 - func-ptr comparison
 - ... ?
 - *The only way to make the standard applicable to real world programs.*

More Rabbit Holes for the Curious

- Relocation details
- Weak linkage
- Versioning (both of libraries and of API)
- Granularity
 - COMDAT, -ffunction-sections, -fdata-sections
- Optimizations
 - Identical Code Folding: /OPT:ICF, -icf=all, -icf=safe
 - Dead Code Elimination: /OPT:REF, -fvtable-gc, --gc-section
- Linker scripts

Resources

- Ulrich Drepper: “How to write shared libraries”
 - <http://library.bagrintsev.me/CPP/dsohowto.pdf>
- Eli Benderski:
 - <https://eli.thegreenplace.net/2011/08/25/load-time-relocation-of-shared-libraries/>
 - <https://eli.thegreenplace.net/2011/11/03/position-independent-code-pic-in-shared-libraries/>
- Rui Ueyama, author of LLD and MOLD linkers:
 - <https://maskray.me/blog/>
- Ian Lance Taylor, author of GOLD:
 - <https://lwn.net/Articles/276782/>
- John Levine, “Linkers and Loaders” book:
 - <https://www.amazon.com/Linkers-Kaufmann-Software-Engineering-Programming/dp/1558604960>
- Michael Kerrisk, Online Training:
 - <https://www.man7.org/training/shlib/index.html>